

# CS 559 Machine Learning

## Linear Classification

Yue Ning  
Department of Computer Science  
Stevens Institute of Technology

# Outline

Generative vs Discriminative Classification

Linear Discriminant Analysis

The Perceptron Algorithm

Naive Bayes Classifier

Model Selection

# Learning Objectives

1. Understand and able to implement Linear Discriminant Analysis;
2. Understand Perceptron and its limitations;

# Classification

Classification task: finding a function  $f$  that classifies examples into a given set of categories  $\{C_1, C_2, \dots, C_k\}$



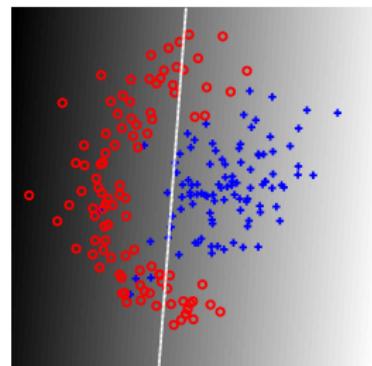
A classification example:



# Binary Classification

Task: Assign each data point to one of two classes.

Examples:

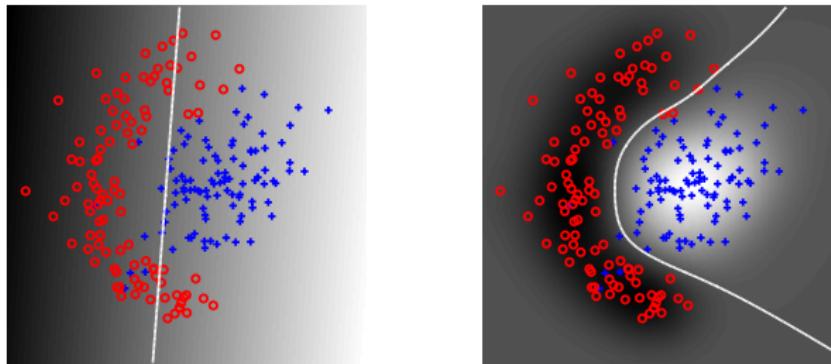


- ▶ Is there a face in this image?
- ▶ Will this neuron spike in response to this stimulus?
- ▶ Based on this brain-scan, does this patient have a given disease or not?
- ▶ Will this customer buy this product or not?
- ▶ Is this person likely to be a democrat/republican?

Notation: we have data

$D = \{(x_1, y_1), \dots, (x_N, y_N)\}$ , with  $y_n = 1$  if  $x_n$  belongs to class 1 and  $y_n = -1$  if  $x_n$  belongs to class  $-1$ .

# Linear Discriminant Functions



Of course, linear algorithms can be used together with **nonlinear feature spaces** or **nonlinear basis functions** in order to solve nonlinear classification problems!

Linear discriminants separate the space by a hyperplane, and the parameters define its normal vector.

- ▶ Decision function:  $f(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + \omega_0$
- ▶ Classification:
  - if  $f(\mathbf{x}) > 0$  say  $\mathbf{x}$  belongs to class 1
  - if  $f(\mathbf{x}) < 0$  say  $\mathbf{x}$  belongs to class -1
- ▶ The decision-surface has equation  $f(\mathbf{x}) = 0$ , and is a hyperplane of dimensionality  $D - 1$ .
- ▶  $\mathbf{w}$  is the normal vector to the hyperplane which separates points into the positive class or negative class.
- ▶  $\omega_0$  determines the location of the decision-surface
- ▶  $|f(\mathbf{x})|$  is proportional to the perpendicular distance to the decision-surface (with factor 1 if  $\|\mathbf{w}\| = 1$ ).

# Linear Discriminant Functions-Geometrical Properties

- ▶ Decision boundary:

$$f(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + \omega_o = 0$$

- ▶ Let  $\mathbf{x}_1, \mathbf{x}_2$  be two points which lie on the decision boundary

$$\begin{aligned} f(\mathbf{x}_1) &= \mathbf{w}^\top \mathbf{x}_1 + \omega_o = 0, f(\mathbf{x}_2) = \mathbf{w}^\top \mathbf{x}_2 + \omega_o = 0 \\ \Rightarrow \mathbf{w}^\top (\mathbf{x}_1 - \mathbf{x}_2) &= 0 \end{aligned}$$

- ▶  $\mathbf{w}$  represents the orthogonal direction to the decision boundary.

# Linear Discriminant Functions-Geometrical Properties

## Cont.

- ▶  $\mathbf{w}^*{}^T = \frac{\mathbf{w}^T}{\|\mathbf{w}\|}$
- ▶  $\mathbf{w}^*{}^T(\mathbf{x} - \mathbf{x}_0)$  is the projection of  $(\mathbf{x} - \mathbf{x}_0)$  onto the  $\mathbf{w}^*$  direction;  $\mathbf{x}_0$  is a point on the decision boundary;
- ▶ Thus,

$$\begin{aligned}\frac{\mathbf{w}^T}{\|\mathbf{w}\|}(\mathbf{x} - \mathbf{x}_0) &= \frac{1}{\|\mathbf{w}\|}(\mathbf{w}^T \mathbf{x} - \mathbf{w}^T \mathbf{x}_0) \\ &= \frac{1}{\|\mathbf{w}\|}(\mathbf{w}^T \mathbf{x} + \omega_0) = \frac{f(\mathbf{x})}{\|\mathbf{w}\|}\end{aligned}$$

when  $\mathbf{x} = \mathbf{0}$ ,  $\frac{f(\mathbf{x})}{\|\mathbf{w}\|} = \frac{\omega_0}{\|\mathbf{w}\|}$

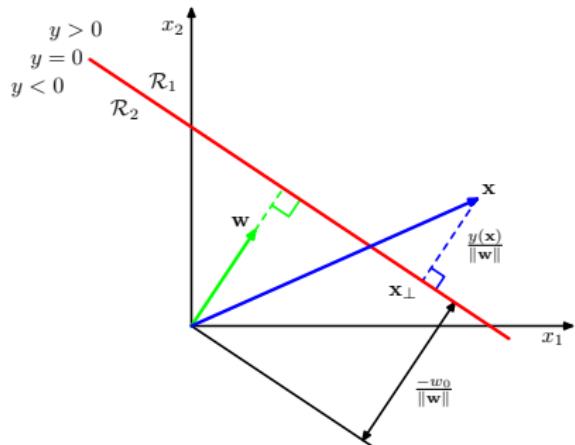
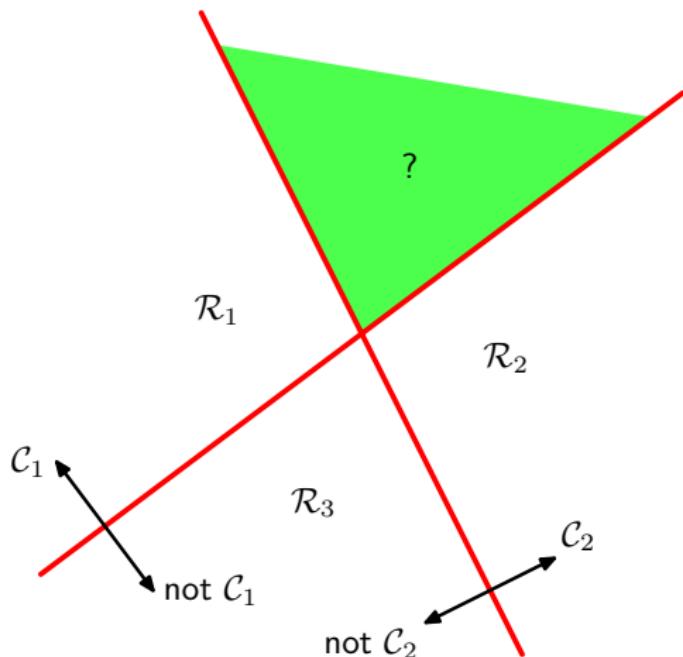


Figure: Signed orthogonal distance of the origin from the decision

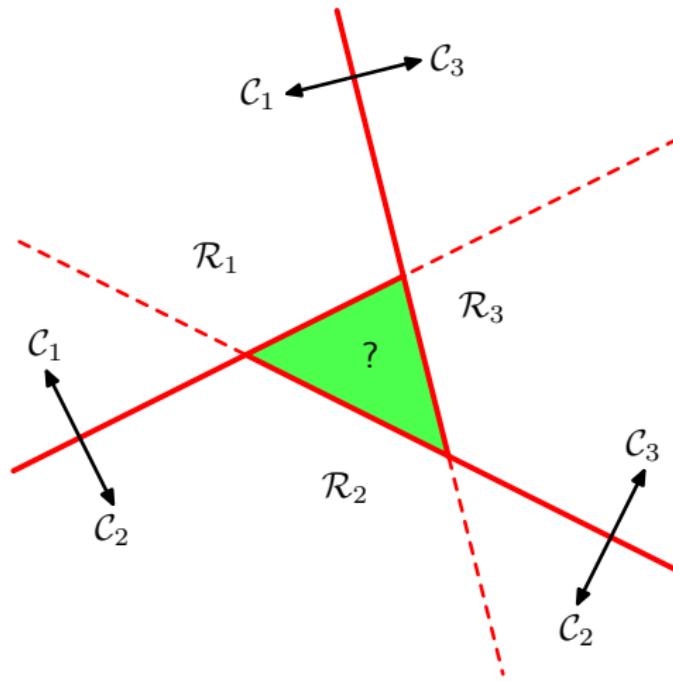
# Linear Discriminant Functions: Multiple classes

one-versus-the-rest:  $K-1$  classifiers each of which solves a two-class problem of separating points of  $C_k$  from points not in that class.



# Linear Discriminant Functions: Multiple classes

one-versus-one:  $\frac{K(K-1)}{2}$  binary discriminant functions, one for every possible pair of classes.

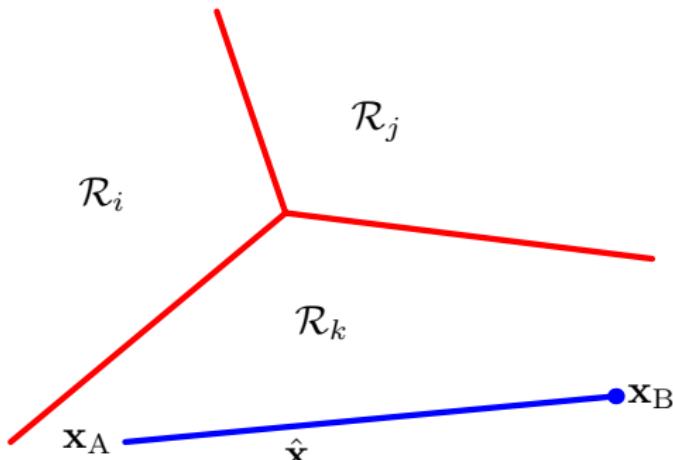


## Linear Discriminant Functions: Multiple classes

- ▶ Solution: consider a single K-class discriminant comprising K linear functions of the form

$$f_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k0}$$

- ▶ Assign a point  $\mathbf{x}$  to class  $C_k$  if  $f_k(\mathbf{x}) > f_j(\mathbf{x}) \forall j \neq k$
- ▶ The decision boundary between class  $C_k$  and class  $C_j$  is given by:  $f_k(\mathbf{x}) = f_j(\mathbf{x}) \Rightarrow (\mathbf{w}_k - \mathbf{w}_j)^T \mathbf{x} + (w_{k0} - w_{j0}) = 0$

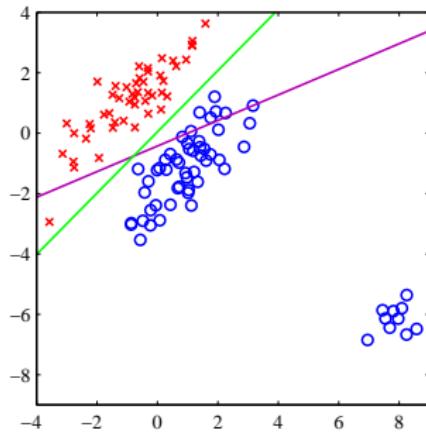
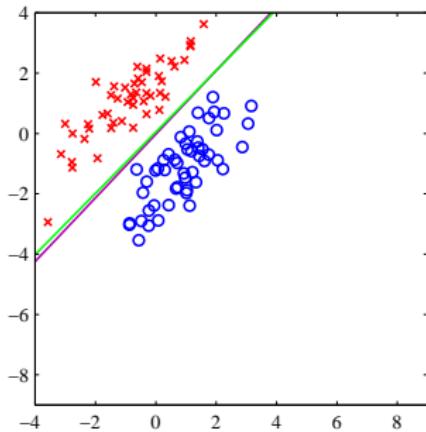


## Multiple algorithms and methods

- ▶ Mis-classification rate  $C(\mathbf{w}) = \frac{1}{N} \sum_n \delta [f(\mathbf{x}_n) = y_n]$  (i.e. average number of errors) difficult to optimize over  $\mathbf{w}$ , and might have multiple solutions.
- ▶ Many algorithms can be derived by replacing  $C$  with another cost function that can be optimized.
- ▶ Linear classification algorithms:
  1. Least-square classification
  2. Fisher's linear Discriminant
  3. Logistic regression
  4. Support Vector Machines
  5. Rosenblatts' perceptron

# Least square classification

- ▶ We have to fit the function  $f(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + \omega_o$  to data.
- ▶ Simply do a linear regression from  $\mathbf{x}$  to  $y$  by minimizing the sum-of-squared errors  $\sum_n (f(\mathbf{x}_n) - y_n)^2$ .
- ▶  $\mathbf{w}_{reg} = (\sum_n \mathbf{x}_n \mathbf{x}_n^\top)^{-1} \sum_n \mathbf{x}_n y_n$
- ▶ Q: In what situations might this be a bad idea?



Bishop PRML Figure 4.4

# Least square classification

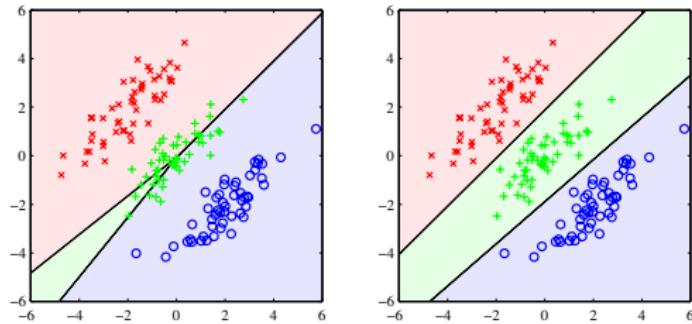


Figure: Left: using a least-squares discriminant; Right: using logistic regression

Bishop PRML Figure 4.5

# Classification via projection

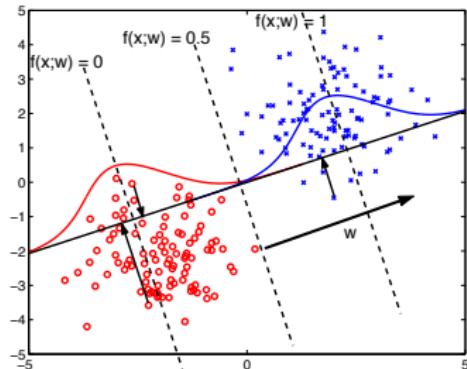
- ▶ A linear function:  $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + \omega_0$  assuming in 2D, projects each point  $\mathbf{x} = [x_1, x_2]^T$  to a line parallel to  $\mathbf{w}$ :

point in $\mathcal{R}^d$	projected point in $\mathcal{R}$
$\mathbf{x}_1$	$z_1 = \mathbf{w}^T \mathbf{x}_1$
$\mathbf{x}_2$	$z_2 = \mathbf{w}^T \mathbf{x}_2$
...	...
$\mathbf{x}_n$	$z_n = \mathbf{w}^T \mathbf{x}_n$

- ▶ We can study how well the projected points  $z_1, \dots, z_n$  are separated across the classes when they are viewed as functions of  $\mathbf{w}$ .

# Classification via projection

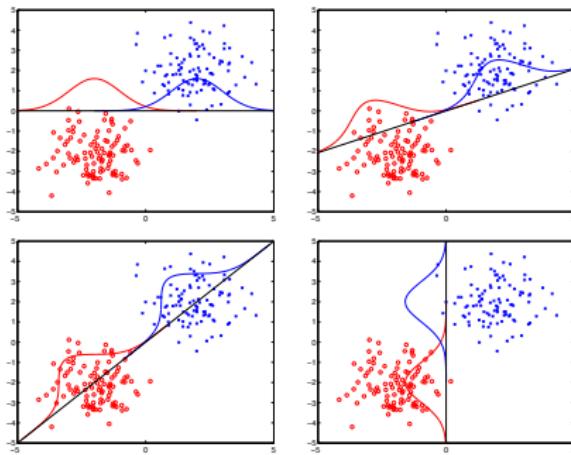
- ▶ A linear function:  $f(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + \omega_0$  assuming in 2D, projects each point  $\mathbf{x} = [x_1, x_2]^\top$  to a line parallel to  $\mathbf{w}$ :



- ▶ We can study how well the projected points  $z_1, \dots, z_n$  viewed as functions of  $\mathbf{w}$  are separated across the classes.

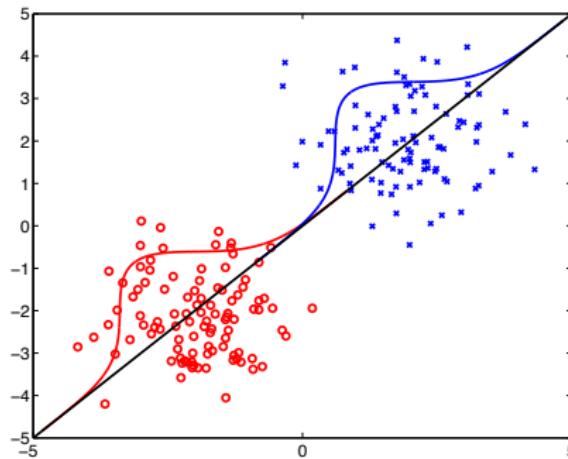
# Classification via projection

- ▶ By varying  $w$  we get different levels of separation between the projected points



# Optimizing the projection

- ▶ We would like to find  $\mathbf{w}$  that somehow maximizes the separation of the projected points across classes.



- ▶ We can quantify the separation (overlap) in terms of means and variances of the resulting 1-dimensional class distributions

# Fisher's linear discriminant

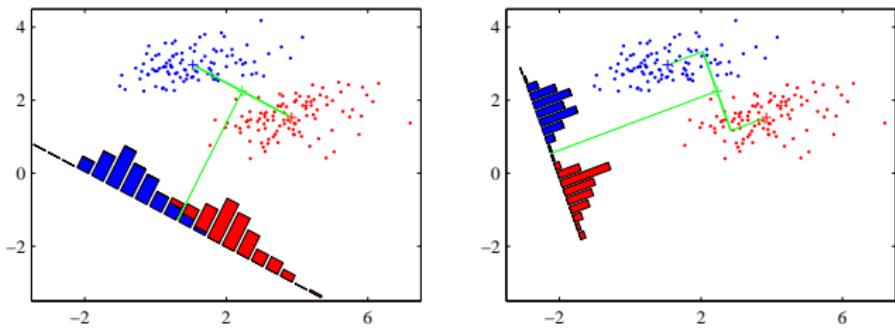
- ▶ One way to view a linear classification model is in terms of dimensionality reduction.
- ▶ Two class case: suppose we project  $\mathbf{x}$  onto one dimension:

$$f = \mathbf{w}^T \mathbf{x}$$

- ▶ Set a threshold  $t$ :

if  $f \leq t$  assign  $C_1$  to  $\mathbf{x}$   
otherwise assign  $C_2$  to  $\mathbf{x}$

# Fisher's linear discriminant



- ▶ Find an orientation along which the projected samples are well separated;
- ▶ This is exactly the goal of linear discriminant analysis (LDA);
- ▶ In other words: we are after the linear projection that best separates the data, i.e. best discriminates data of different classes.

# Fisher's linear discriminant

- ▶ Two classes:  $\{C_+, C_-\}$
- ▶  $N_+$  samples of class  $C_+$
- ▶  $N_-$  samples of class  $C_-$
- ▶ Consider  $\mathbf{w} \in \mathbb{R}^d$  with  $\|\mathbf{w}\| = 1$
- ▶ Then:  $\mathbf{w}^T \mathbf{x}$  is the projection of  $\mathbf{x}$  along the direction of  $\mathbf{w}$ .
- ▶ We want the projections  $\mathbf{w}^T \mathbf{x}$  where  $\mathbf{x} \in C_+$  separated from the projections  $\mathbf{w}^T \mathbf{x}$  where  $\mathbf{x} \in C_-$

# Fisher's linear discriminant

- ▶ A measure of the separation between the projected points is the difference of the sample means:
  - Sample mean of class  $C_+$ :

$$\mathbf{m}_+ = \frac{1}{N_+} \sum_{\mathbf{x} \in C_+} \mathbf{x}$$

- Sample mean for the projected points:

$$m_+ = \frac{1}{N_+} \sum_{\mathbf{x} \in C_+} \mathbf{w}^T \mathbf{x} = \mathbf{w}^T \mathbf{m}_+$$

$$\Rightarrow |m_+ - m_-| = \mathbf{w}^T (\mathbf{m}_+ - \mathbf{m}_-)$$

- ▶ We wish to make the above difference as large as we can. In addition, ...

# Fisher's linear discriminant

- ▶ To obtain good separation of the projected data, we really want the difference between the means to be large relative to some measure of the standard deviation of each class:
  - Scatter of the projected samples of class  $C_+$ :

$$s_+^2 = \sum_{\mathbf{x} \in C_+} (\mathbf{w}^T \mathbf{x} - m_+)^2$$

- Total within-class scatter of the projected samples:

$$s_+^2 + s_-^2$$

- Fisher linear discriminant analysis:

$$\arg \max_{\mathbf{w}} \frac{|m_+ - m_-|^2}{s_+^2 + s_-^2}$$

# Fisher's linear discriminant

- ▶  $J(\mathbf{w}) = \frac{|m_+ - m_-|^2}{s_+^2 + s_-^2}$
- ▶ To obtain  $J(\mathbf{w})$  as an explicit function of  $\mathbf{w}$ , we define the following matrices:

$$S_+ = \sum_{\mathbf{x} \in C_+} (\mathbf{x} - \mathbf{m}_+)(\mathbf{x} - \mathbf{m}_+)^T$$

Within-class scatter matrix:

$$S_w = S_+ + S_-$$

- ▶ Then:

$$\begin{aligned}s_+^2 &= \sum_{\mathbf{x} \in C_+} (\mathbf{w}^T \mathbf{x} - m_+)^2 = \sum_{\mathbf{x} \in C_+} (\mathbf{w}^T \mathbf{x} - \mathbf{w}^T \mathbf{m}_+)^2 \\ &= \sum_{\mathbf{x} \in C_+} \mathbf{w}^T (\mathbf{x} - \mathbf{m}_+) (\mathbf{x} - \mathbf{m}_+)^T \mathbf{w} = \mathbf{w}^T S_+ \mathbf{w}\end{aligned}$$

## Fisher's linear discriminant

- ▶ So,  $s_+^2 = \mathbf{w}^T S_+ \mathbf{w}$  and  $s_-^2 = \mathbf{w}^T S_- \mathbf{w}$
- ▶ Thus,

$$\begin{aligned}s_+^2 + s_-^2 &= \mathbf{w}^T S_+ \mathbf{w} + \mathbf{w}^T S_- \mathbf{w} \\&= \mathbf{w}^T (S_+ + S_-) \mathbf{w} \\&= \mathbf{w}^T S_w \mathbf{w}\end{aligned}$$

- ▶ Similarly:

$$\begin{aligned}(m_+ - m_-)^2 &= (\mathbf{w}^T \mathbf{m}_+ - \mathbf{w}^T \mathbf{m}_-)^2 \\&= \mathbf{w}^T (\mathbf{m}_+ - \mathbf{m}_-) (\mathbf{m}_+ - \mathbf{m}_-)^T \mathbf{w} \\&= \mathbf{w}^T S_B \mathbf{w}\end{aligned}$$

where  $S_B = (\mathbf{m}_+ - \mathbf{m}_-) (\mathbf{m}_+ - \mathbf{m}_-)^T$  (**Between-class scatter matrix**)

# Fisher's linear discriminant

- We have obtained:

$$J(\mathbf{w}) = \frac{|m_+ - m_-|^2}{s_+^2 + s_-^2} = \frac{\mathbf{w}^T S_B \mathbf{w}}{\mathbf{w}^T S_W \mathbf{w}}$$

- $J(\mathbf{w})$  is maximized when  $(\mathbf{w}^T S_B \mathbf{w}) S_W \mathbf{w} = (\mathbf{w}^T S_W \mathbf{w}) S_B \mathbf{w}$
- We observe that:

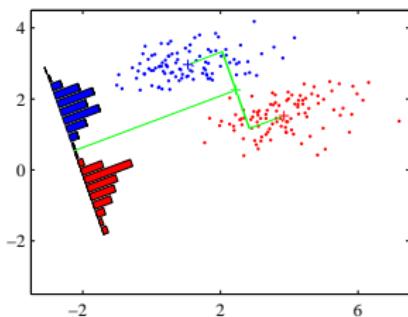
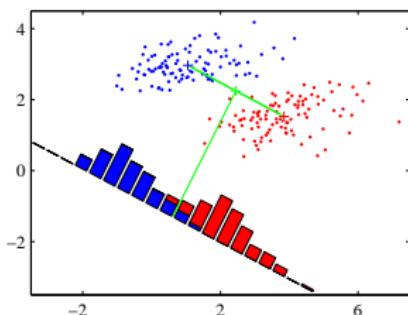
$$S_B \mathbf{w} = (\mathbf{m}_+ - \mathbf{m}_-) (\mathbf{m}_+ - \mathbf{m}_-)^T \mathbf{w}$$

where  $(\mathbf{m}_+ - \mathbf{m}_-)^T \mathbf{w}$  is a scalar, always in the direction of  $(\mathbf{m}_+ - \mathbf{m}_-)$

- Solution:

$$\mathbf{w} = S_W^{-1} (\mathbf{m}_+ - \mathbf{m}_-)$$

# Fisher's linear discriminant: Summary



- ▶  $\mathbf{m}_+ = \frac{1}{N_+} \sum_{n \in C_+} \mathbf{x}_n$
- ▶  $\mathbf{m}_- = \frac{1}{N_-} \sum_{n \in C_-} \mathbf{x}_n$
- ▶ Maximize projection-distance of class means  $\mathbf{w}_{simple} \propto \mathbf{m}_+ - \mathbf{m}_-$
- ▶ Maximizing distance between means ignores that the projected variances might also be big.
- ▶ Fix: Maximize the ratio of between-class variance to within-class variance ('signal to noise'). Fisher criterion

$$J_{\mathbf{w}} = \frac{(m_+ - m_-)^2}{s_+^2 + s_-^2} \quad (1)$$

$$\mathbf{w}_{Lda} = S_W^{-1}(\mathbf{m}_+ - \mathbf{m}_-)$$

Bishop PRML Figure 4.6

# Fisher's linear discriminant

- ▶ Gives the linear function with the maximum ratio of between-class scatter to within-class scatter.
- ▶ The problem, e.g. classification, has been reduced from a  $d$ -dimensional problem to a more manageable one-dimensional problem.
- ▶ Optimal for multivariate normal class conditional densities.

## Fisher's linear discriminant: Multi-Class

- ▶ The analysis can be extended to multiple classes.
- ▶  $S_W = \sum_{k=1}^K \sum_{\mathbf{x}_i \in C_k} (\mathbf{x}_i - \mathbf{m}_k)(\mathbf{x}_i - \mathbf{m}_k)^T$
- ▶  $S_B = \sum_{k=1}^K m_k(\mathbf{m}_k - \mathbf{m})(\mathbf{m}_k - \mathbf{m})^T$  where  $\mathbf{m}$  is the global mean;  $m_k$  is the number of examples in class k
- ▶ Solve:  $S_B \mathbf{v} = \lambda S_W \mathbf{v}$  the generalized eigenvalue problem
- ▶ At most K-1 distinct solution eigenvalues
- ▶ The optimal projection matrix  $V$  to a subspace of dimension  $k$  is given by the eigenvectors corresponding to the largest  $k$  eigenvalues

# Fisher's linear discriminant

- ▶ LDA is a linear technique for **dimensionality** reduction: it projects the data along directions that can be expressed as linear combination of the input features.
- ▶ The “appropriate” transformation depends on the data and on the task we want to perform on the data. Note that LDA uses class labels.
- ▶ **Non-linear** extensions of LDA exist (e.g., generalized LDA).

# The Perceptron Algorithm (Frank Rosenblatt, 1957)

- ▶ First learning algorithm for neural networks.
- ▶ Originally introduced for character classification, where each character is represented as an image;
- ▶ Total input to output node:

$$\sum_j w_j x_j$$

- ▶ Output unit performs the function (activation function):

$$H(x) = \begin{cases} 1 & \text{if } x \geq 0 \\ 0 & \text{if } x < 0 \end{cases}$$

# Perceptron: Learning Algorithm

- ▶ **Goal:** compute a mapping from inputs to the outputs.
- ▶ **Example:** two class character recognition problem.
  - Training set: set of images representing either the character 'a' or the character 'b' (supervised learning);
  - Learning task: learn the weights so that when a new unlabelled image comes in, the network can predict its label.
  - Setting:  $d$  input units (intensity level of a pixel), 1 output unit.

# Perceptron: Learning Algorithm

The algorithm proceeds as follows:

- ▶ Initial random setting of weights;
- ▶ The input is a random sequence  $\{\mathbf{x}_k\}$
- ▶ For each element of class  $C_1$ , if output = 1 (correct), do nothing; otherwise, update weights;
- ▶ For each element of class  $C_2$ , if output = 0 (correct), do nothing; otherwise, update weights;

# Perceptron: Learning Algorithm

- ▶ More formally:  $\mathbf{x} = (x_1, x_2, \dots, x_d)^T$ ,  $\mathbf{w} = (w_1, w_2, \dots, w_d)^T$
- ▶  $\theta$ : Threshold of the output unit
- ▶ Unit output:  $\mathbf{w}^T \mathbf{x} = w_1 x_1 + w_2 x_2 + \dots + w_d x_d$
- ▶ Output class 1 if  $\mathbf{w}^T \mathbf{x} - \theta \geq 0$
- ▶ To eliminate the explicit dependence on  $\theta$ : Output class 1 if:  $\mathbf{w}^T \mathbf{x} \geq 0$

# Perceptron: Learning Algorithm

- ▶ We want to learn values of the weights so that the perceptron correctly discriminate elements of  $C_1$  from elements of  $C_2$
- ▶ Given  $x$  in input, if  $x$  is classified correctly, weights are unchanged, otherwise:

$$w = \begin{cases} w + x & \text{if an element of class } C_1 \text{ was classified as in } C_2 \\ w - x & \text{if an element of class } C_2 \text{ was classified as in } C_1 \end{cases}$$

# Perceptron: Learning Algorithm

- ▶ **1<sup>st</sup> case:**  $\mathbf{x} \in C_1$  and was classified in  $C_2$ . The correct answer is 1, which corresponds to:  $\mathbf{w}^T \mathbf{x} \geq 0$ , we have  $\mathbf{w}^T \mathbf{x} < 0$ . We want to get closer to the correct answer:  $\mathbf{w}^T \mathbf{x} < \mathbf{w}'^T \mathbf{x}$ .

$$\mathbf{w}^T \mathbf{x} < \mathbf{w}'^T \mathbf{x}, \text{ iff } \mathbf{w}^T \mathbf{x} < (\mathbf{w} + \mathbf{x})^T \mathbf{x}$$

$$(\mathbf{w} + \mathbf{x})^T \mathbf{x} = \mathbf{w}^T \mathbf{x} + \mathbf{x}^T \mathbf{x} = \mathbf{w}^T \mathbf{x} + \|\mathbf{x}\|^2$$

because  $\|\mathbf{x}\|^2 > 0$ , the condition is verified.

- ▶ **2<sup>nd</sup> case:**  $\mathbf{x} \in C_2$  and was classified in  $C_1$ . The correct answer is 0, which corresponds to:  $\mathbf{w}^T \mathbf{x} < 0$ , we have  $\mathbf{w}^T \mathbf{x} \geq 0$ . We want to get closer to the correct answer:  $\mathbf{w}^T \mathbf{x} > \mathbf{w}'^T \mathbf{x}$ .

$$\mathbf{w}^T \mathbf{x} > \mathbf{w}'^T \mathbf{x}, \text{ iff } \mathbf{w}^T \mathbf{x} < (\mathbf{w} - \mathbf{x})^T \mathbf{x}$$

$$(\mathbf{w} - \mathbf{x})^T \mathbf{x} = \mathbf{w}^T \mathbf{x} - \mathbf{x}^T \mathbf{x} = \mathbf{w}^T \mathbf{x} - \|\mathbf{x}\|^2$$

because  $\|\mathbf{x}\|^2 > 0$ , the condition is verified.

# Perceptron: Learning Algorithm

In summary:

- ▶ A random sequence  $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_k$  is generated such that  $x_i \in C_1 \cup C_2$
- ▶ If  $\mathbf{x}_k$  is correctly classified, then  $\mathbf{w}_{k+1} = \mathbf{w}_k$  otherwise:

$$\mathbf{w}_{k+1} = \begin{cases} \mathbf{w}_k + \mathbf{x}_k & \text{if } \mathbf{x}_k \in C_1 \\ \mathbf{w}_k - \mathbf{x}_k & \text{if } \mathbf{x}_k \in C_2 \end{cases}$$

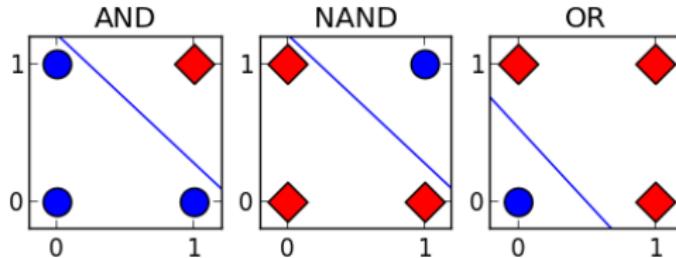
- ▶ Convergence theorem: regardless of the initial choice of weights, if the two classes are linearly separable, there exists  $\mathbf{w}$  such that:

$$\begin{cases} \mathbf{w}^T \mathbf{x} \geq 0 & \text{if } \mathbf{x}_k \in C_1 \\ \mathbf{w}^T \mathbf{x} < 0 & \text{if } \mathbf{x}_k \in C_2 \end{cases}$$

then the learning rule will find such solution after a finite number of steps.

# Representational Power of Perceptrons

- ▶ Marvin Minsky and Seymour Papert, “Perceptrons” 1969:  
The perceptron can solve only problems with linearly separable classes
- ▶ Examples of linearly separable Boolean functions:



# Representational Power of Perceptrons

- ▶ Examples of linearly separable Boolean functions:

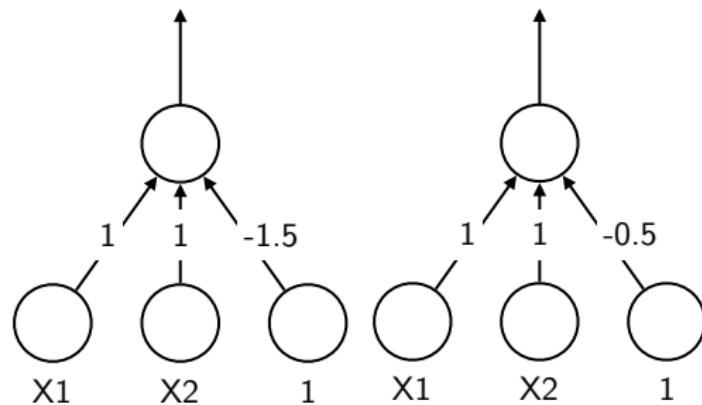
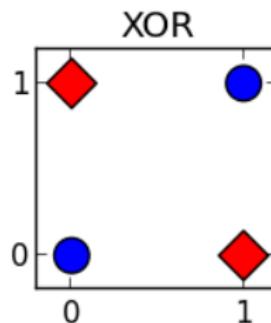


Figure: Left: AND; Right: OR

# Representational Power of Perceptrons

- ▶ Examples of a non linearly separable Boolean function:



- ▶ The EX-OR function cannot be computed by a perceptron.

## Naive Bayes: not (necessarily) a Bayesian method;

- ▶ A and B are independent iff  $p(A, B) = p(A)p(B)$
- ▶ A and B are conditionally independent given C iff  
 $p(A, B|C) = p(A|C)p(B|C)$

# Naive Bayes: Assumption

- ▶ Assume dimensions of  $\mathbf{x}$  are conditionally independent given  $y$ .
- ▶ Example, bag of words:  
$$p(\text{"Stevens"}, \text{"Institute"}, \text{"Technology"} | y) =$$
$$p(\text{"Stevens"} | y)p(\text{"Institute"} | y)p(\text{"Technology"} | y)$$
- ▶ Optimizing:

$$\begin{aligned} f(x) &= \arg \max_y p(y|x) \\ &= \arg \max_y p(x|y)p(y)/p(x) \\ &= \arg \max_y p(x|y)p(y) \\ &= \arg \max_y p(y) \prod_j p(x_j|y) \end{aligned}$$

# Naive Bayes: Solution

- ▶  $p(y) \leftarrow \frac{\# \text{ examples where } Y=y}{(\# \text{ examples})}$
- ▶  $p(X_j = x_j | y) \leftarrow \frac{\# \text{ ex where } Y=y \text{ and } X_j=x_j}{(\# \text{ ex where } Y=y)}$
- ▶ Learning by counting!

## Gaussian naive Bayes: Continuous data

- ▶  $p(y) := \frac{\# \text{ examples where } Y=y}{(\# \text{ examples})}$
- ▶  $p(X_j = v|y) \leftarrow \frac{1}{\sqrt{2\pi\sigma_k^2}} \exp\left\{-\frac{(v-\mu_k)^2}{2\sigma_k^2}\right\}$
- ▶  $\mu_k$  and  $\sigma_k$  are determined from the training data set.
- ▶ Learning by counting!

# Gaussian naive Bayes: example (from Wikipedia)

Training data set:

Sex	height	weight	foot size
male	6	180	12
male	5.92	1990	11
male	5.58	170	12
male	5.92	165	10
female	5	100	6
female	5.5	150	8
female	5.42	140	7
female	5.75	150	9

Mean and variance

Sex	mean-height	var-height	mean-weight	var-weight	mean-footsize	var-footsize
male	5.855	$3.5 * 10^{-2}$	176.25	$1.2292 * 10^2$	11.25	$9.1667 * 10^{-1}$
female	5.4175	$9.7225 * 10^{-2}$	132.5	$5.5833 * 10^2$	7.5	1.6667

# Gaussian naive Bayes: example (from Wikipedia)

Training mean and variance

Sex	mean-height	var-height	mean-weight	var-weight	mean-footsize	var-footsize
male	5.855	$3.5 * 10^{-2}$	176.25	$1.2292 * 10^2$	11.25	$9.1667 * 10^{-1}$
female	5.4175	$9.7225 * 10^{-2}$	132.5	$5.5833 * 10^2$	7.5	1.6667

Testing:

Sex	height	weight	foot size
?	6	130	8

$$p(m|x) \approx p(m)p(\text{height}|m)p(\text{weight}|m)p(\text{footsize}|m) = 6.1984 * 10^{-9}$$

$$p(f|x) \approx p(f)p(\text{height}|f)p(\text{weight}|f)p(\text{footsize}|f) = 5.3778 * 10^{-4}$$

# What is Model Selection?

Given a set of models  $\mathcal{M} = \{M_1, M_2, \dots, M_R\}$ , choose the model that is expected to do the best on the **test data**.  $\mathcal{M}$  may consist of:

- ▶ Same learning model with **different complexities** or **hyperparameters**.
  - Nonlinear regression: polynomials with different degrees
  - K-Nearest Neighbors: Different choices of K
  - Decision Trees: Different choices of the number of levels/leaves
  - SVM: Different choices of the misclassification penalty
  - Regularized models: Different choices of the regularization parameter
  - Kernel based methods: Different choices of kernels ...and almost any learning problem
- ▶ Different **learning models** (e.g. SVM, kNN, DT, etc)

**Note:** usually considered in supervised learning but unsupervised learning faces this issue too.

## Held-out Data

- ▶ Set aside a fraction (10-20%) of the training data.
- ▶ This part becomes our held-out data (validation/development)
- ▶ **Remember:** Held-out data is NOT the test data
- ▶ Train each model using the remaining training data
- ▶ Evaluate error on the held-out data
- ▶ Choose the model with the smallest held-out error
- ▶ **Problems:**
  - wastes training data
  - if there was an unfortunate split (can be alleviated by repeated random subsampling)

# Cross-Validation

## K-fold Cross-Validation on $N$ training examples

- ▶ Create  **$K$  equal sized partitions** of the training data
- ▶ Each partition has  $N/K$  examples
- ▶ Train using  $K - 1$  partitions, validate on the remaining partition
- ▶ Repeat the same  $K$  times, each with a different validation partition
- ▶ Choose the model with the **smallest average validation error**
- ▶ Usually  $K$  is chosen as 10

# Leave-One-Out (LOO) Cross-Validation

Special case of K-fold Cross-Validation when  $K = N$

- ▶ Each partition is now **an example**
- ▶ Train using  $N - 1$  examples, validate on the remaining example
- ▶ Repeat the same  $N$  times, each with a different validation example
- ▶ Choose the model with the **smallest average validation error**
- ▶ **can be expensive** for large  $N$ . Typically used when  $N$  is small

# Random Subsampling Cross-Validation

- ▶ Randomly subsample a fixed fraction  $\alpha N (0 < \alpha < 1)$  of examples; call it the validation set
- ▶ Training using the rest of the examples, measure error on the validation set
- ▶ Repeat  $K$  times, each with a different randomly chosen validation set
- ▶ Choose the model with the smallest average validation error
- ▶ Usually  $\alpha$  is chose as 0.1,  $K$  as 10

# Bootstrapping

- ▶ Given a set of  $N$  examples
- ▶ Idea: Sample  $N$  elements from this set with **replacement**  
(already sampled elements can be picked again)
- ▶ Use this new set as the training data
- ▶ The set of examples not selected as the validation data
- ▶ For large  $N$ , training data consists of about only **63% unique** examples
- ▶ Expected model error:

$$e = 0.632 \times e_{\text{test}} + 0.368e_{\text{training}}$$

- ▶ This can break down if we overfit and  $e_{\text{training}} = 0$

Bradley Efron & Robert Tibshirani. *Improvements on Cross-Validation: The 63+ Bootstrap Method*

# Information Criteria based methods

- ▶ Akaike Information Criteria (AIC)

$$\text{AIC} = 2k - 2 \log(\mathcal{L})$$

- ▶ Bayesian Information Criteria (BIC)

$$\text{BIC} = k \log(N) - 2 \log(\mathcal{L})$$

- ▶ k: # of model parameters
- ▶ N: # of data examples
- ▶  $\mathcal{L}$ : maximum value of the model likelihood
- ▶ Applicable for probabilistic models
- ▶ AIC/BIC penalize model complexity

# Feature Selection

Selecting a useful subset from all the features. Why?

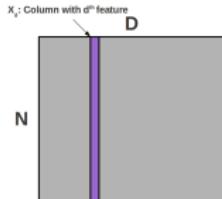
- ▶ Some algorithms **scale (computationally) poorly** with increased dimension
- ▶ **Irrelevant** features can confuse some algorithms
- ▶ **Redundant** features adversely affect regularization
- ▶ Removal of features can **increase (relative) margin** (and generalization)
- ▶ Reduces data set and resulting model size
- ▶ Note: Feature Selection is different from Feature Extraction
  - The latter transforms original features to get a small set of new features
  - More on feature extraction when we cover **Dimensionality Reduction**

# Feature Selection Methods

- ▶ Methods agnostic to the learning algorithm
  - Preprocessing based methods
    - E.g., remove a binary feature if it's ON in very few or most examples
  - Filter Feature Selection methods
    - Use some ranking criteria to rank features
    - Select the top ranking features
- ▶ Wrapper Methods (keep the learning algorithm in the loop)
  - Requires repeated runs of the learning algorithm with different set of features
  - Can be computationally expensive

# Filter Feature Selection

- ▶ Uses heuristics but is much faster than wrapper methods



- ▶ **Correlation Criteria:** Rank features in order of their correlation with the labels

$$R(X_d, \mathbf{y}) = \frac{\text{cov}(X_d, \mathbf{y})}{\sqrt{\text{var}(X_d)\text{var}(\mathbf{y})}}$$

- ▶ **Mutual Information Criteria:**

$$MI(X_d, \mathbf{y}) = \sum_{X_d \in \{0,1\}} \sum_{y \in \{-1,+1\}} P(X_d, \mathbf{y}) \log \frac{P(X_d, \mathbf{y})}{P(X_d)P(y)}$$

- high mutual information means high relevance of that feature
- Note: these probabilities can be easily estimated from the data

# Wrapper Methods

- ▶ Forward Search
  - Start with no features
  - Greedily include the most relevant feature
  - Stop when selected the desired number of features
- ▶ Backward Search
  - Start with all features
  - Greedily remove the least relevant feature
  - Stop when selected the desired number of features
- ▶ Inclusion/Removal criteria uses cross-validation

## Acknowledgements

Slides adapted from Dr. Bert Huang's *Machine Learning* at Virginia Tech, Dr. Tommi S. Jaakkola's *Introduction to Machine Learning* at MIT, and Dr. Piyush Rai's *Machine Learning* at Utah.