

CS 559 Machine Learning

Boosting

Yue Ning
Department of Computer Science
Stevens Institute of Technology

Plan for today

Boosting

Reduce Bias and Decrease Variance

- ▶ Bagging reduces variance by averaging
- ▶ Bagging has little effect on bias
- ▶ Can we average and reduce bias?
- ▶ Yes: **Boosting**

Example: Email Spam

How would you classify an email as SPAM or not? using following criteria. If:

1. Email has only one image file (promotional image), **SPAM**
2. Email has only link(s), **SPAM**
3. Email body consist of sentence like “You won a prize money of \$ xxxxxx”, **SPAM**
4. Email from our official domain “stevens.edu” , **Not a SPAM**
5. Email from known source, **Not a SPAM**

The Boosting Approach

- ▶ devise computer program for deriving rough rules
- ▶ apply procedure to subset of examples
- ▶ obtain a simple rule
- ▶ apply to 2nd subset of examples
- ▶ obtain a 2nd rule
- ▶ repeat T times

Key details

- ▶ How to choose examples on each round?
 - concentrate on “hardest” examples (those most often misclassified by previous rule)
- ▶ How to combine the rules into single prediction rule?
 - take (weighted) majority vote of rules

Boosting

- ▶ **boosting** = general method of converting rough rules into highly accurate prediction rule
- ▶ **technically**
 - assume given “weak” learning algorithm that can consistently find classifiers at least slightly better than random, say, accuracy $\geq 55\%$
 - given sufficient data, a boosting algorithm can provably construct single classifier with very high accuracy say, 99%

A formal description of boosting

- ▶ Given **training set** $(x_1, y_1), \dots, (x_N, y_N)$
- ▶ $y_i \in \{-1, +1\}$ correct labels of instance $x_i \in X$
- ▶ for $t = 1, \dots, T$:
 - construct weight distribution $u^{(t)}$ on $\{1, \dots, N\}$
 - find weak classifier:

$$f_t : X \rightarrow \{-1, +1\}$$

with error ϵ_t on $u^{(t)}$:

$$\epsilon = P_{i \sim u^{(t)}}[f_t(x_i) \neq y_i]$$

- ▶ Output final/combined classifier H_{final}

The Idea of AdaBoost

- ▶ Training $f_2(x)$ on the new training set that fails $f_1(x)$
- ▶ How to find a new training set that fails $f_1(x)$? ϵ_1 : the error rate of $f_1(x)$ on its training data

$$\epsilon_1 = \frac{\sum_n u_n^{(1)} \delta(f_1(x_n) \neq y_n)}{Z_1}$$

where $Z_1 = \sum_n u_n^{(1)}$, $\epsilon_1 < 0.5$

- ▶ Changing the example weights from $u_n^{(1)}$ to $u_n^{(2)}$ such that:
 $\frac{\sum_n u_n^{(2)} \delta(f_1(x_n) \neq y_n)}{Z_2} = 0.5$ The performance of f_1 for new weights would be random
- ▶ Training $f_2(x)$ based on the new weights $u_n^{(2)}$

AdaBoost

► constructing $u^{(t)}$

- $u_i^{(1)} = 1/n$
- given $u^{(t)}$ and f_t :

$$\begin{aligned} u_i^{(t+1)} &= \frac{u_i^{(t)}}{Z^t} \times \begin{cases} e^{-\alpha_t} & \text{if } y_i = f_t(x_i) \\ e^{\alpha_t} & \text{if } y_i \neq f_t(x_i) \end{cases} \\ &= \frac{u_i^{(t)}}{Z^t} \times \exp(-\alpha_t y_i f_t(x_i)) \end{aligned}$$

where


Z^t = the normalization factor

$$\alpha_t = \frac{1}{2} \ln\left(\frac{1 - \epsilon_t}{\epsilon_t}\right) > 0$$

► final classifier



$$H_{\text{final}}(x) = \text{sign}\left(\sum_t \alpha_t f_t(x)\right)$$

Toy example



weak classifiers = vertical or horizontal half-planes


Round 1




$$\epsilon_1 = 0.30$$

$$\alpha_1 = 0.42$$


Round 2



Round 3



Final Classifier



Voted combination of classifiers

- ▶ The general problem here is to try to combine many simple “weak” classifiers into a single “strong” classifier
- ▶ We consider voted combinations of simple binary component classifiers

$$f_T(x) = \alpha_1 f(\mathbf{x}; \theta_1) + \dots + \alpha_T f(\mathbf{x}; \theta_T)$$

where θ is the model parameter and the (non-negative) votes α_i can be used to emphasize component classifiers that are more reliable than others.

The AdaBoost algorithm

1. Set $u_i^{(0)} = 1/n$ for $i = 1, \dots, n$
2. At the m^{th} iteration we find (any) classifier $f(\mathbf{x}; \hat{\theta}_m)$ for which the weighted classification error ϵ_m

$$\epsilon_m = 0.5 - \frac{1}{2} \left(\sum_{i=1}^n u_i^{(m-1)} y_i f(\mathbf{x}_i; \hat{\theta}_m) \right)$$

is better than chance

3. The new component is assigned votes based on its error
(smaller error rate, larger weight for voting)

$$\hat{\alpha}_m = 0.5 \log \frac{1 - \epsilon_m}{\epsilon_m}$$

4. The weights are updated according to (Z_m is chosen so that the new weights $u_i^{(m)}$ sum to one):

$$u_i^{(m)} = \frac{1}{Z_m} \cdot u_i^{(m-1)} \exp(-y_i \hat{\alpha}_m f(\mathbf{x}_i; \hat{\theta}_m))$$

Acknowledgements

Slides adapted from Dr. David Sontag's *Machine Learning* at MIT.