

# CS 559 Machine Learning

## Linear Regression

Yue Ning  
Department of Computer Science  
Stevens Institute of Technology

# Plan for today

Linear Regression

Gradient Descent

Learning Rate

Features

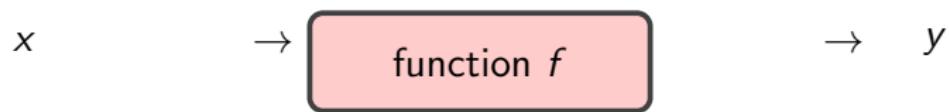
Beyond Linear Regression Models

Bias and Variance

The Curse of Dimensionality

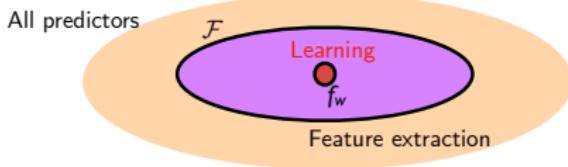
Maximum Likelihood Estimator

# Regression Task



# The Learning Problem

- ▶ **Hypothesis class:** we consider some restricted set  $\mathcal{F}$  of mappings  $f : \mathcal{X} \rightarrow \mathcal{Y}$  from input data to output.



- ▶ **Estimation:** on the basis of a set of training examples with their labels,  $\{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$ , we find an estimate  $\hat{f} \in \mathcal{F}$
- ▶ **Evaluation:** we measure how well  $\hat{f}$  generalizes to yet unseen examples, i.e., whether  $\hat{f}(\mathbf{x}_{\text{new}})$  agrees with  $y_{\text{new}}$

## Loss

A loss function  $\text{Loss}(\mathbf{x}, \mathbf{y}, \mathbf{w})$  quantifies how wrong you would be if you used  $\mathbf{w}$  to make a prediction on  $\mathbf{x}$  when the correct output is  $\mathbf{y}$ . It is the object we want to minimize.

- ▶ Loss is a function of the parameters  $\mathbf{w}$  and we can try to minimize it directly.
- ▶ We reduce the estimation problem to a minimization problem.
- ▶ Valid for any parameterized class of mapping from examples to predictions.
- ▶ Valid when the predictions are discrete labels or real valued as long as the loss is defined properly.

# Linear Regression

## Prediction

$$f_w(\mathbf{x}) = \hat{y} = \mathbf{w}^T \mathbf{x}$$

## Residual

The residual is the amount by which prediction overshoots the target.

## Squared loss

$$L(\mathbf{x}, y, \mathbf{w}) = (\underbrace{f_w(\mathbf{x}) - y}_{\text{residual}})^2$$

Example:

$$\mathbf{w} = [2, -1], \mathbf{x} = [2, 0], y = -1, L(\mathbf{x}, y, \mathbf{w}) = 25$$

# Linear Regression

Prediction:  $f_w(\mathbf{x}) = \hat{y} = \mathbf{w}^T \mathbf{x} = w_0 + w_1 x$

$$\text{Loss: } L(\mathbf{x}, y, \mathbf{w}) = \underbrace{(f_{\mathbf{w}}(\mathbf{x}) - y)^2}_{\text{residual}}$$

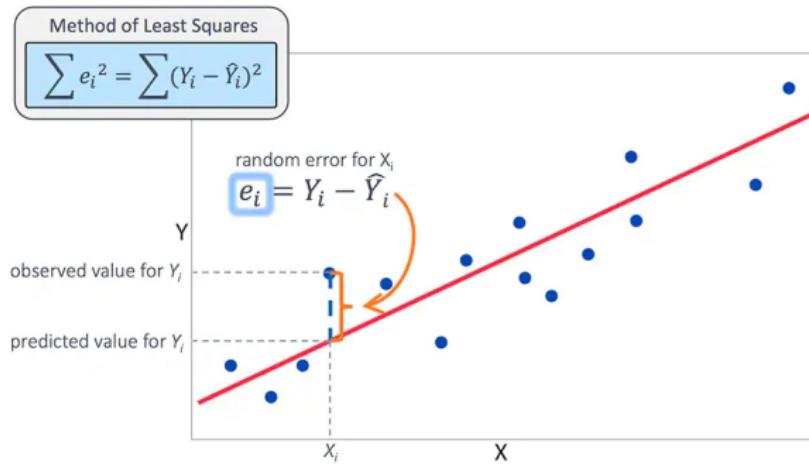


Figure: Fitting one-dimensional linear regression

# Loss Minimization Framework

Loss is easy to minimize if there is only one example.

## Minimize Total Training Loss

$$\min_{\mathbf{w} \in \mathbb{R}^d} \text{TrainLoss}(\mathbf{w}) = \underbrace{\frac{1}{N} \sum_{(x,y) \in D_{\text{train}}} L(\mathbf{x}, y, \mathbf{w})}_{\text{empirical loss}}$$

Key: learn a  $\mathbf{w}$  to make global tradeoffs ( $N$  is the number of examples in  $D_{\text{train}}$ ).

## Training and Test performance: sampling

- ▶ Assume each training and test example-label pair  $(\mathbf{x}, y)$  is drawn independently at random from the same but unknown population of examples and labels.
- ▶ We can represent this population as a joint probability distribution  $P(\mathbf{x}, y)$  so that each training/test example is a sample from this distribution  $(\mathbf{x}_i, y_i) \sim P$ .

$$\text{Empirical (training) loss} = \frac{1}{N} \sum_i \text{Loss}(y_i, f_w(\mathbf{x}))$$

$$\text{Expected (test) loss} = E_{(\mathbf{x}, y) \sim P} \{ \text{Loss}(y_i, f_w(\mathbf{x})) \}$$

- ▶ The training loss based on a few sampled examples and labels serves as a proxy for the test performance measured over the whole population.

# Regression Loss Functions

## Squared loss

$$L(\mathbf{x}, y, \mathbf{w})_2 = (f_{\mathbf{w}}(\mathbf{x}) - y)^2$$

## Absolute loss

$$L(\mathbf{x}, y, \mathbf{w})_1 = |f_{\mathbf{w}}(\mathbf{x}) - y|$$

# Which loss to use?

Example:

$$D_{\text{train}} = \{(1, 0), (1, 2), (1, 1000)\}$$

- ▶ For least square regression (L2):  $\mathbf{w}$  that minimizes training loss is mean  $\mathbf{y}$

$$\frac{\partial L_2}{\partial f_w} = -\frac{2}{N} \sum_{i=1}^N (y_i - f_w) = 0 \rightarrow f_w = \frac{1}{N} \sum_{i=1}^N y_i$$

- ▶ Mean: tries to accommodate every example, popular

# Which loss to use?

Example:

$$D_{\text{train}} = \{(1, 0), (1, 2), (1, 1000)\}$$

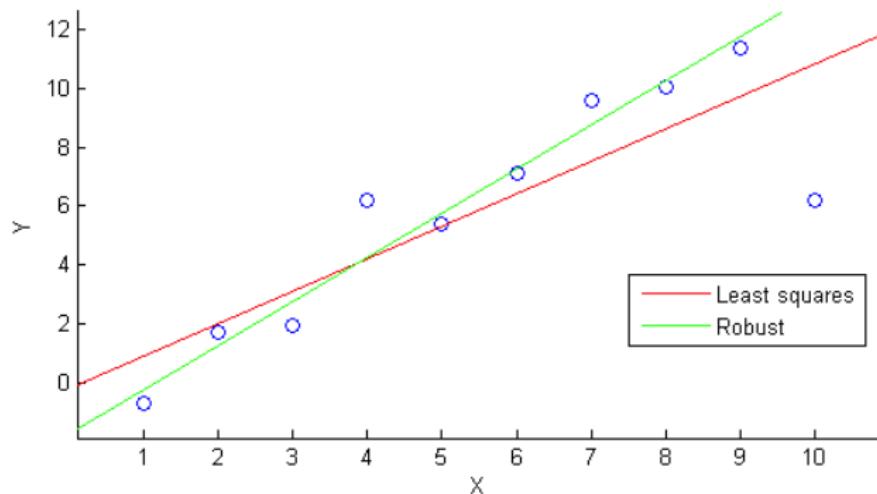
- ▶ For least absolute deviation regression (L1):  $\mathbf{w}$  that minimizes training loss is median  $y$

$$\frac{\partial L_1}{\partial f_w} = -\frac{1}{N} \sum_{i=1}^N \text{sgn}(y_i - f_w)$$

(The derivative equals to 0 when there is the same number of positive and negative terms among  $y_i - f_w$ , which means  $f_w$  should be the median of  $y_i$ )

- ▶ Median: more robust with outliers.

# Comparison



Least squares:

$$Y = -0.188327 + 1.10351*X$$

RMS error = 2.21375

Robust:

$$Y = -1.77278 + 1.50415*X$$

RMS error = 1.43663

Figure: Least squares vs. robust regression; Demo from Mathworks

# Squared loss: Closed Form Solution (Normal Equation)

- ▶ L2 loss:

$$\text{Squared loss} = \frac{1}{N} \sum_{(x,y) \in D_{\text{train}}} (\mathbf{w}^\top \mathbf{x} - y)^2$$

(Set the gradient w.r.t  $\mathbf{w}$  to be zero?)

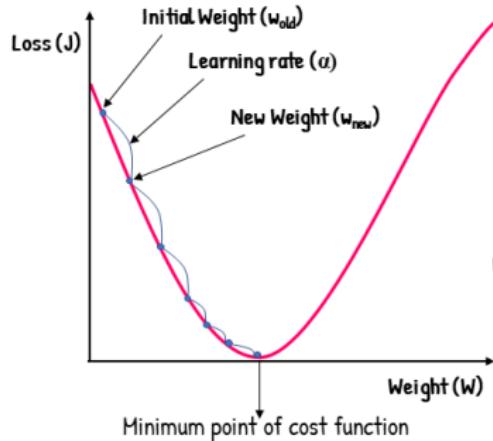
- ▶ Solution:

$$\mathbf{w}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

Reading: Lineare regression lecture at Cornell CS  
Linear regression on Wikipedia

# Another solution for optimization: Gradient Descent

## Gradient Descent



$$w_{\text{new}} = w_{\text{old}} - \alpha \frac{\delta J}{\delta w}$$

Figure: From Understanding Gradient Descent Algorithm and Its Role in Linear Regression

# Linear regression using gradient descent

## Objective function

$$\text{TrainLoss}(\mathbf{w}) = L(\mathbf{w}) = \frac{1}{N} \sum_{(\mathbf{x}, y) \in D_{\text{train}}} (\mathbf{w}^T \mathbf{x} - y)^2$$

## Gradient with respect to $\mathbf{w}$ (use chain rule)

$$\nabla_{\mathbf{w}} \text{TrainLoss}(\mathbf{w}) = \nabla_{\mathbf{w}} L(\mathbf{w}) = \frac{1}{N} \sum_{(\mathbf{x}, y) \in D_{\text{train}}} 2(\underbrace{\mathbf{w}^T \mathbf{x}}_{\text{prediction}} - \underbrace{y}_{\text{target}}) \mathbf{x}$$

# Gradient descent is slow

## Gradient

$$\nabla_{\mathbf{w}} L(\mathbf{w}) = \frac{1}{N} \sum_{(x,y) \in D_{\text{train}}} 2(\underbrace{\mathbf{w}^T \mathbf{x}}_{\text{prediction}} - \underbrace{y}_{\text{target}}) \mathbf{x}$$

## Gradient descent update

$$\mathbf{w} \leftarrow \mathbf{w} - \underbrace{\eta}_{\text{step size}} \underbrace{\nabla_{\mathbf{w}} L(\mathbf{w})}_{\text{gradient}}$$

Each iteration requires going over all training examples – expensive when have lots of data

# Stochastic Gradient Descent

## Gradient Descent (GD)

$$\mathbf{w} \leftarrow \mathbf{w} - \underbrace{\eta}_{\text{step size}} \underbrace{\nabla_{\mathbf{w}} L(\mathbf{w})}_{\text{gradient}}$$

## Stochastic Gradient Descent (SGD)

For each  $(\mathbf{x}, y) \in D_{train}$

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} L_{(\mathbf{x}, y)}(\mathbf{w})$$

**Key:** It's not about quality, it's about quantity.

# (Batch) Gradient Descent

---

```
1: procedure BATCH GRADIENT DESCENT
2:   Initialize w with random values
3:   for i in range(epochs) do
4:      $g^{(i)}(\mathbf{w}) = \text{evaluate\_gradient}(\text{TrainLoss}, \text{data}, \mathbf{w})$ 
5:      $\mathbf{w} = \mathbf{w} - \text{learning\_rate} * g^{(i)}(\mathbf{w})$ 
6:   Return w from the break point
```

---

- ▶ Can be very slow.
- ▶ Intractable for datasets that don't fit in memory.
- ▶ Doesn't allow us to update our model online, i.e. with new examples on-the-fly.
- ▶ Guaranteed to converge to the global minimum for convex error surfaces and to a local minimum for non-convex surfaces.

# Stochastic Gradient Descent

---

```
1: procedure STOCHASTIC GRADIENT DESCENT
2:   Initialize w with random values
3:   for i in range(epochs) do
4:     np.random.shuffle(data)
5:     for example ∈ data do
6:        $g^{(i)}(\mathbf{w}) = \text{evaluate\_gradient}(\text{loss, example, } \mathbf{w})$ 
7:        $\mathbf{w} = \mathbf{w} - \text{learning\_rate} * g^{(i)}(\mathbf{w})$ 
8:     Check stopping criteria
9:   return w from the break point
```

---

- ▶ Allow for online update with new examples.
- ▶ With a high variance that cause the objective function to fluctuate heavily.

# Mini-batch Gradient Descent

---

```
1: procedure MINI-BATCH GRADIENT DESCENT
2:   Initialize w with random values
3:   for  $i$  in range(epochs) do
4:     np.random.shuffle(data)
5:     for batch  $\in$  get_batches(data, batch_size=50) do
6:        $g^{(i)}(\mathbf{w}) = \text{evaluate\_gradient}(\text{loss}, \text{batch}, \mathbf{w})$ 
7:        $\mathbf{w} = \mathbf{w} - \text{learning\_rate} * g^{(i)}(\mathbf{w})$ 
8:     Check stopping criteria
9:   return w from the break point
```

---

- ▶ Reduces the variance of the parameter updates, which can lead to more stable convergence;
- ▶ Can make use of highly optimized matrix optimizations common to state-of-the-art deep learning libraries that make computing the gradient w.r.t. a mini-batch very efficient.

# Learning Rate

Gradient Update:

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} \text{Loss}(x, y, \mathbf{w})$$

Question: what should the step size (learning rate) be?

Strategies:

For each  $(x, y) \in D_{train}$

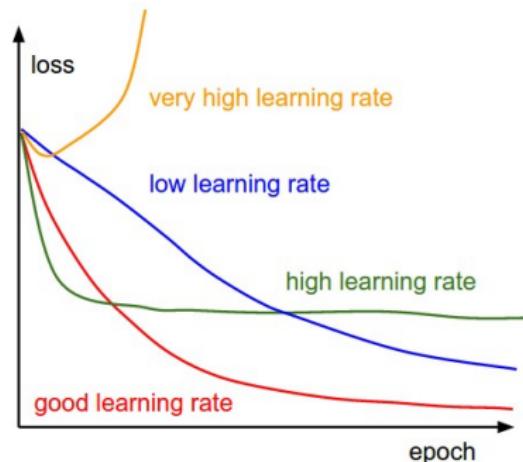
- ▶ Constant:  $\eta = 0.1$
- ▶ Decreasing:  $\eta = \frac{1}{\sqrt{\# \text{ updates made so far}}}$

# Learning Rate

Set the learning rate  $\eta$  carefully:

$$\mathbf{w}^{(i)} \leftarrow \mathbf{w}^{(i-1)} - \eta \nabla L(\mathbf{w}^{(i-1)})$$

If there are more than 3 parameters, it is hard to visualize the loss w.r.t the parameters.  
But you can visualize the loss w.r.t the # of parameter updates (epoch).



# Adaptive Learning Rates

- ▶ Popular & Simple Idea: Reduce the learning rate by some factor every few epochs.
  - At the beginning, we are far from the destination, so we use larger learning rate
  - After several epochs, we are close to the destination, so we reduce the learning rate
  - e.g.  $\frac{1}{t}$  decay:  $\eta^{(t)} = \frac{\eta}{\sqrt{t+1}}$
- ▶ Learning rate cannot be one-size-fits-all
  - Giving different parameters different learning rates
  - AdaGrad! [paper in Journal of Machine Learning Research]

Reading: An overview of gradient descent optimization algorithms

# Adagrad

- We have

$$\eta^{(t)} = \frac{\eta}{\sqrt{t+1}}, \quad g_i^{(t)} = \frac{\partial L(x, y, \mathbf{w}_i^{(t)})}{\partial w_i^{(t)}}$$

- Vanilla Gradient descent

$$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \eta^{(t)} g^{(t)}$$

- Adagrad: Divide the learning rate of each parameter by the root mean square of its previous derivatives

$$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \frac{\eta^{(t)}}{\sigma^{(t)}} g^{(t)}$$

where  $\sigma^{(t)}$  is the root mean square of the previous derivatives of parameter  $w$

# Adagrad

$$w^{(1)} \leftarrow w^{(0)} - \frac{\eta^{(0)}}{\sigma^{(0)}} g^{(0)}$$

$$\sigma^{(0)} = \sqrt{(g^{(0)})^2 + \epsilon}$$

$$w^{(2)} \leftarrow w^{(1)} - \frac{\eta^{(1)}}{\sigma^{(1)}} g^{(1)}$$

$$\sigma^{(1)} = \sqrt{\frac{1}{2}[(g^{(0)})^2 + (g^{(1)})^2] + \epsilon}$$

$$w^{(3)} \leftarrow w^{(2)} - \frac{\eta^{(2)}}{\sigma^{(2)}} g^{(2)}$$

$$\sigma^{(2)} = \sqrt{\frac{1}{3}[(g^{(0)})^2 + (g^{(1)})^2 + (g^{(2)})^2] + \epsilon}$$

...

...

$$w^{(t+1)} \leftarrow w^{(t)} - \frac{\eta^{(t)}}{\sigma^{(t)}} g^{(t)}$$

$$\sigma^{(t)} = \sqrt{\frac{1}{t+1} \sum_{i=1}^t (g^{(i)})^2 + \epsilon}$$

$\epsilon$  is a smoothing term that avoids division by zero (usually on the order of  $1e-8$ )

# Summary so far

- ▶ Linear predictors:  
 $f_w(x)$  based on score  $\mathbf{w}^T \mathbf{x}$
- ▶ Training loss minimization: learning as optimization:

$$\min_w L(w)$$

- ▶ Stochastic gradient descent: optimization algorithm:

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} L_{(\mathbf{x}, y)}(\mathbf{w})$$

## Two components

Score:

$$\mathbf{w}^T \mathbf{x} = \sum_{i=1}^d w_i x_i$$

- ▶ Previous: learning  $\mathbf{w}$  via optimization
- ▶ Next: feature extraction and vector representation based on prior domain knowledge

## Example: Data

**Example:** specifies that  $y$  is the ground-truth output for  $x$

$$(x, y)$$

**Training data:** a set of examples

$$D_{\text{train}} = [(\text{"...located in Hoboken, NJ..."}, 1.4M), \\ (\text{"...Fairfax, VA..."}, 0.7M), \\ (\text{}, \dots)]$$

# Feature Extraction

- ▶ Example task: given a house, predict house price  $y$
- ▶ Question: what properties of  $x$  might be relevant for predicting  $y$ ?
- ▶ Feature extractor: Given input  $x$ , output a set of (**feature name**, **feature value**) pairs.

average neighborhood house price:	\$1.1M
average school score:	5.6
distance to train stations:	3.4(m)
parking garage:	1
size:	1200(sqrt)

# Feature Vector

Mathematically, feature vector doesn't need feature names:

$$\begin{bmatrix} \text{average neighborhood house price: } \$1.1M \\ \text{average school score: } 5.6 \\ \text{distance to train stations: } 3.4(\text{m}) \\ \text{parking garage: } 1 \\ \text{size: } 1200(\text{sqrt}) \end{bmatrix} \rightarrow \begin{bmatrix} 1.1 \\ 5.6 \\ 3.4 \\ 1 \\ 1200 \end{bmatrix}$$

## Feature vector

For input  $X$ , its feature vector is:

$$\mathbf{x} = x_1, \dots, x_d$$

$\mathbf{x}$  as a point in a high-dimensional space.

# Linear Predictors

Weight vector  $\mathbf{w} \in \mathbb{R}^d$

$$\begin{bmatrix} \text{avg\_n\_price} & 1.8 \\ \text{avg\_school} & 1.4 \\ \text{dist\_train:} & 0.8 \\ \text{parking:} & 1.5 \\ \text{size:} & 1.6 \end{bmatrix}$$

Feature vector  $\mathbf{x} \in \mathbb{R}^d$

$$\begin{bmatrix} \text{avg\_n\_price} & \$1.1 \\ \text{avg\_school} & 5.6 \\ \text{dist\_train:} & 3.4 \\ \text{parking:} & 1 \\ \text{size:} & 1200 \end{bmatrix}$$

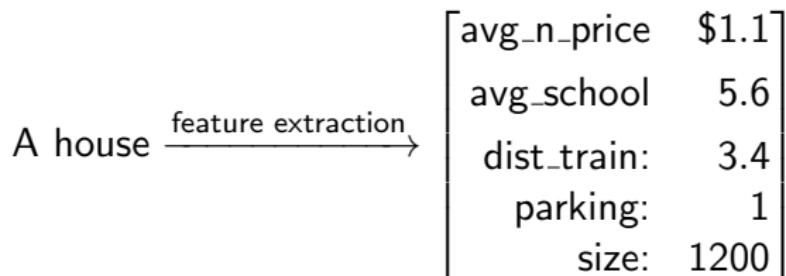
## Score

weighted combination of features

$$\mathbf{w}^T \mathbf{x} = \sum_{i=1}^d w_i x_i$$

# Organization of Features

Task: predict the price of a house



Which features to construct? need prior knowledge

# Feature Templates

A feature template is a group of features all computed in a similar way. Input:

a house located in.....

Some feature templates:

- ▶ Average school score is \_\_
- ▶ Distance to train stations less than \_\_
- ▶ Located in \_\_

# Sparsity in feature vectors

Feature template: located in ...

A house located in ...  $\xrightarrow{\text{feature extraction}}$

located_in_nyc	1
located_in_DC	0
located_in_seattle	0
located_in_pittsburg	0
...	...
located_in_sf	0

Inefficient to represent all the zeros...

# Feature vector representation

avg_n_price	\$1.1
avg_school	5.6
dist_train:	3.4
parking:	1
size:	1200

Array representation: (good for dense features):

[1.1, 5.6, 3.4, 1, 1200]

Map representation: (good for sparse features):

{ "avg\_n\_price" : 1.1, "avg\_school" : 5.6 }

# Example: Vector Representation

Linear functions:

$$\mathbf{x} = [x]$$

$$\mathcal{F}_1 = \{x \rightarrow w_1 x + w_2 x^2 : w_1 \in \mathbb{R}, w_2 = 0\}$$

Quadratic functions:

$$\mathbf{x} = [x, x^2]$$

$$\mathcal{F}_2 = \{x \rightarrow w_1 x + w_2 x^2 : w_1 \in \mathbb{R}, w_2 \in \mathbb{R}\}$$

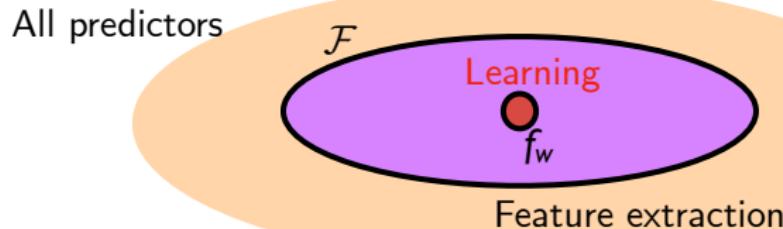
Quadratic functions are supersets of linear functions

# Expressivity

## Hypothesis class

A hypothesis class is the set of possible predictors with a fixed  $\mathbf{x}$  and a varying  $\mathbf{w}$

$$\mathcal{F} = \{f_{\mathbf{w}} : \mathbf{w} \in \mathbb{R}^d\}$$



# Features in linear predictors

## Example: medical diagnosis

- ▶ Input features: height, weight, body temperature, blood pressure, etc.
- ▶ Output: real value.

Three issues: (non-linear in original measurements)

- ▶ Non-monotonicity
- ▶ Saturation
- ▶ Interactions between features

# Non-monotonicity (1)

- ▶ Features, Try(1):

$$\mathbf{x} = [1, \text{temperature}(x)]$$

- ▶ Output:

health  $y \in \mathbb{R}$

- ▶ Problem: favor extremes; true relationship is non-monotonic

## Non-monotonicity (2)

- ▶ Features, Try(2): transform inputs

$$\mathbf{x} = [1, (\text{temperature}(x) - 37)^2]$$

- ▶ Output:

health  $y \in \mathbb{R}$

- ▶ Disadvantage: requires manually-specified domain knowledge

## Non-monotonicity (3)

- ▶ Features, Try(3): transform inputs

$$\mathbf{x} = [1, \text{temperature}(x), \text{temperature}(x)^2]$$

- ▶ Output:

health  $y \in \mathbb{R}$

- ▶ General rule: features should be simple building blocks to be pieced together

# Saturation (1)

Example: product recommendation

- ▶ Input: product information  $x$
- ▶ Output: relevance  $y$

The number of people who bought  $x$

$$\text{Identity: } \mathbf{x} = N(x)$$

Problem: is 1000 people really 10 times more relevant than 100 people? Not quite...

## Saturation (2)

- ▶ The number of people who bought  $x$

$$\text{Identity: } \mathbf{x} = N(x)$$

- ▶ Log

$$\mathbf{x} = \log N(x)$$

- ▶ Discretization

$$\mathbf{x} = [\mathbb{I}(1 < N(x) \leq 10), \mathbb{I}(10 < N(x) \leq 100), \dots]$$

# Interaction between features (1)

Example: predicting health

- ▶ Input: patient information  $x$
- ▶ Output: health  $y$

Try (1):

$$\mathbf{x} = [\text{height}(x), \text{weight}(x)]$$

Problem: cannot capture relationship between height and weight.

## Interaction between features (2)

Try (2):

$$x = (52 + 1.9(\text{height}(x) - 60) - \text{weight}(x))^2$$

**Solution:** define features that combine inputs.

**Disadvantage:** requires manually-specified domain knowledge.

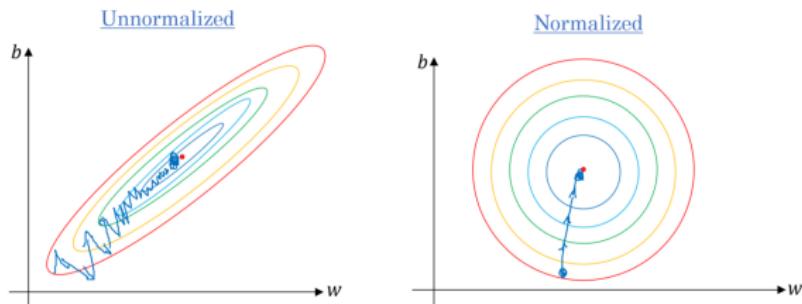
## Interaction between features (3)

Try (3):

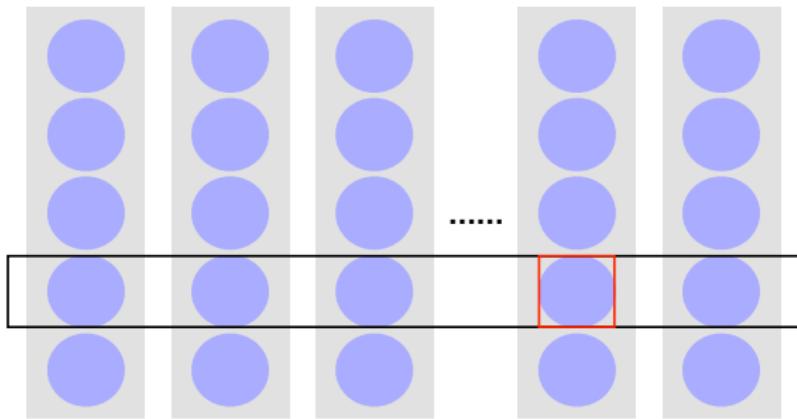
$$\mathbf{x} = [1, \text{height}(x), \text{weight}(x), \text{height}(x)^2, \text{weight}(x)^2, \text{height}(x)\text{weight}(x)]$$

**Solution:** add features involving multiple measurements.

# Feature Scaling



# Feature Scaling



Assuming we have  $n$  ( $r = 1, \dots, n$ ) examples. For each dimension  $i$ , the mean is  $m_i$ , and the standard deviation is  $\sigma_i$ . Scaling features:

$$x_i^r \leftarrow \frac{x_i^r - m_i}{\sigma_i}$$

# Linear in what?

Prediction driven by score:

$$\mathbf{w}^T \mathbf{x}$$

- ▶ Linear in  $\mathbf{w}$ ? Yes
- ▶ Linear in  $\mathbf{x}$ ? Yes
- ▶ Linear in each input feature? No!

## Non-linearity

- ▶ Predictors can be expressive **non-linear** functions and decision boundaries of  $\mathbf{x}$
- ▶ Score is **linear** function of  $\mathbf{w}$ , which permits efficient learning

## Summary so far

- ▶ Goal: define features so that the hypothesis class contains good predictors
- ▶ Pay attention to non-linearity in  $x$ : non-monotonicity, saturation, interaction between features
- ▶ Suggested approach: define features  $\mathbf{x}$  to be building blocks (e.g., monomials)
- ▶ Linear prediction: actually very powerful!

# Closed-form solution

- ▶ Objective:

$$\min L = \min \sum_{n=1}^N \{y_n - \mathbf{w}^T \mathbf{x}_n\}^2$$

- ▶ Transform into matrix operation ( $\mathbf{y}_{N \times 1}$ ,  $\mathbf{X}_{N \times d}$ ,  $\mathbf{w}_{d \times 1}$ ):

$$\min(\mathbf{y} - \mathbf{X}\mathbf{w})^T(\mathbf{y} - \mathbf{X}\mathbf{w})$$

- ▶ Differentiating w.r.t  $\mathbf{w}$ :

$$\mathbf{X}^T(\mathbf{y} - \mathbf{X}\mathbf{w}) = 0$$

- ▶ Setting the gradient to zero, we get a closed form of estimates:

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

where  $\mathbf{X}$  is an  $N \times d$  matrix. **The solution is a linear function of the outputs  $\mathbf{y}$**

- ▶ The resulting prediction errors  $\epsilon_i = y_i - f(\mathbf{x}_i)$  are uncorrelated with any linear function of the inputs  $\mathbf{x}$ .
- ▶ Easily extend to non-linear functions of the inputs.

# Beyond Linear Regression

- ▶ Example extension:  $m^{th}$  order polynomial regression on one-dimensional input where  $f : \mathcal{R} \rightarrow \mathcal{R}$  is given by:

$$f(\mathbf{w}) = w_0 + w_1x + \dots + w_{m-1}x^{m-1} + w_mx^m$$

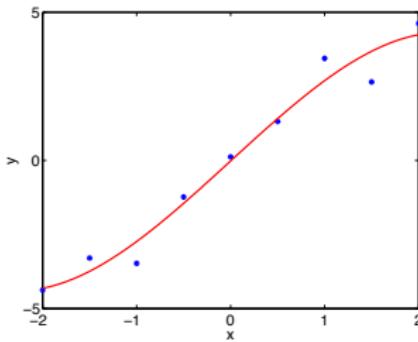
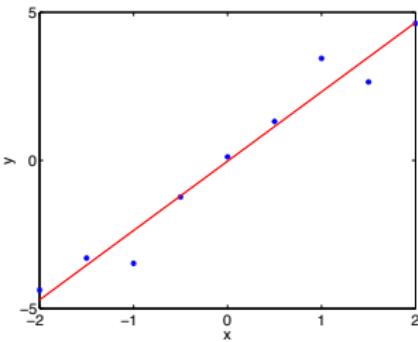
- ▶ Solution as before:

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

where:

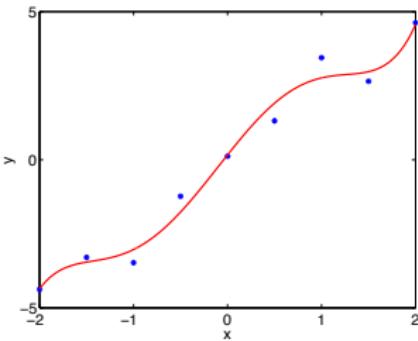
$$\mathbf{w} = \begin{pmatrix} w_0 \\ w_1 \\ \dots \\ w_m \end{pmatrix}, \quad \mathbf{X} = \begin{pmatrix} 1 & x_1 & x_1^2 & \dots & x_1^m \\ 1 & x_2 & x_2^2 & \dots & x_2^m \\ \dots & \dots & \dots & \dots & \dots \\ 1 & x_N & x_N^2 & \dots & x_N^m \end{pmatrix}$$

# Polynomial Regression

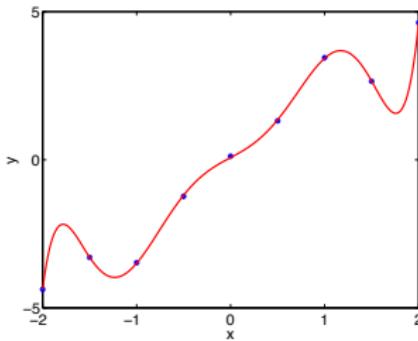


degree = 1

degree = 3



degree = 5



degree = 7

# Underfitting and Overfitting

- ▶ Underfitting:
  - Training an ML algorithm on training data → 0.24 accuracy
  - Applying the ML algorithm on test data → 0.25 accuracy
- ▶ Overfitting:
  - Training an ML algorithm on training data → 0.95 accuracy
  - Applying the ML algorithm on test data → 0.27 accuracy (huge generalization error!)

# Underfitting and Overfitting

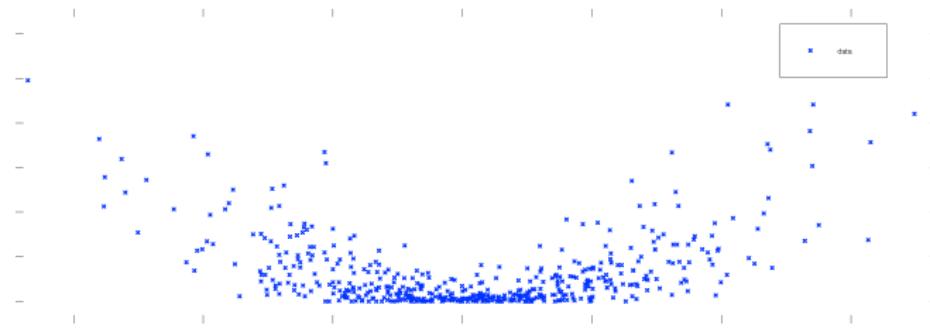


Figure: Regression data

# Underfitting and Overfitting

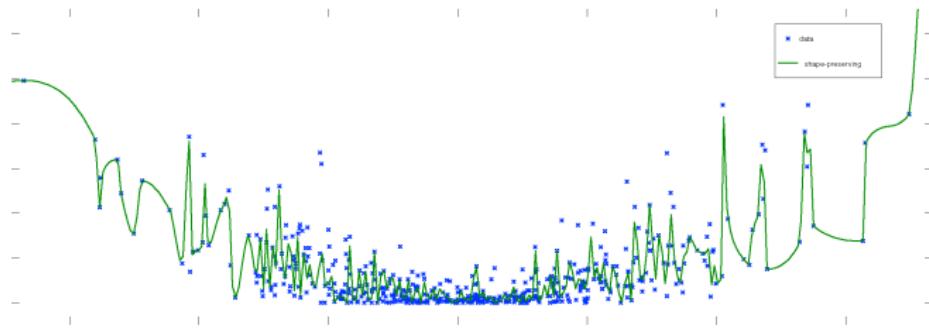


Figure: Overfitting

# Underfitting and Overfitting

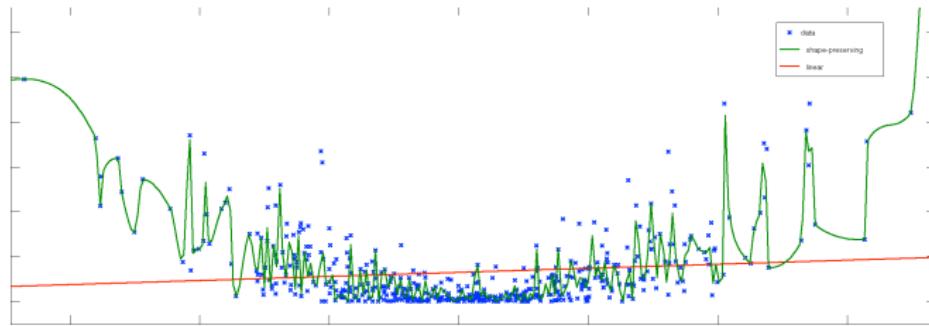


Figure: underfitting

# Underfitting and Overfitting

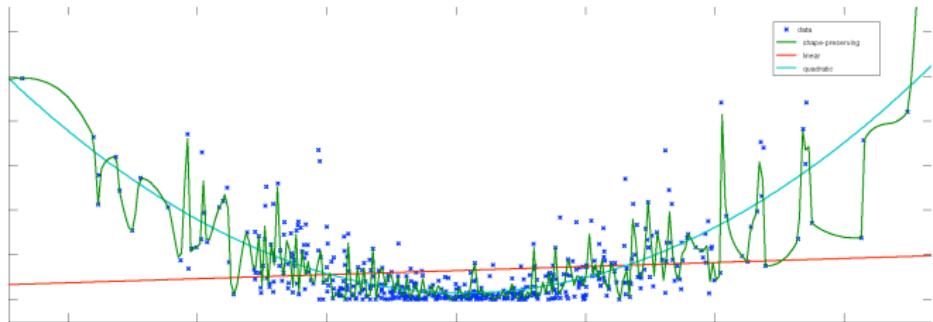


Figure: Overfitting, underfitting, and generalization

# Underfitting and Overfitting

- ▶ Underfitting:
  - Low dimensional
  - Heavily regularized
  - Bad modeling assumption
- ▶ Overfitting:
  - High dimensional or non-parametric
  - Weakly regularized
  - Not enough data

## Summary

- ▶ Training models too complex can cause overfitting
- ▶ Training models too simple (or wrong) can cause underfitting

# Bias and Variance

- ▶ A training set  $(x_1, \dots, x_n)$  and  $y$  values
- ▶ True function  $y = f(x) + \epsilon$  ( $\epsilon$  has mean 0 and variance  $\sigma^2$ ).
- ▶  $\hat{f}(x)$  that approximates the true function  $f(x)$
- ▶ Expected squared prediction error at a point  $x$  is:

$$E[(y - \hat{f}(x))^2] = \underbrace{(E[\hat{f}(x)] - f(x))^2}_{\text{bias}^2} + \underbrace{(E[\hat{f}(x)^2] - E^2[\hat{f}(x)])}_{\text{Variance}} + \sigma^2$$

where

$$\text{Bias}[\hat{f}(x)] = E[\hat{f}(x)] - f(x), \text{Var}[\hat{f}(x)] = E[\hat{f}(x)^2] - E^2[\hat{f}(x)]$$

- ▶ Bias: error from incorrect modeling assumption
- ▶ Variance: error from random noise

# Bias and Variance

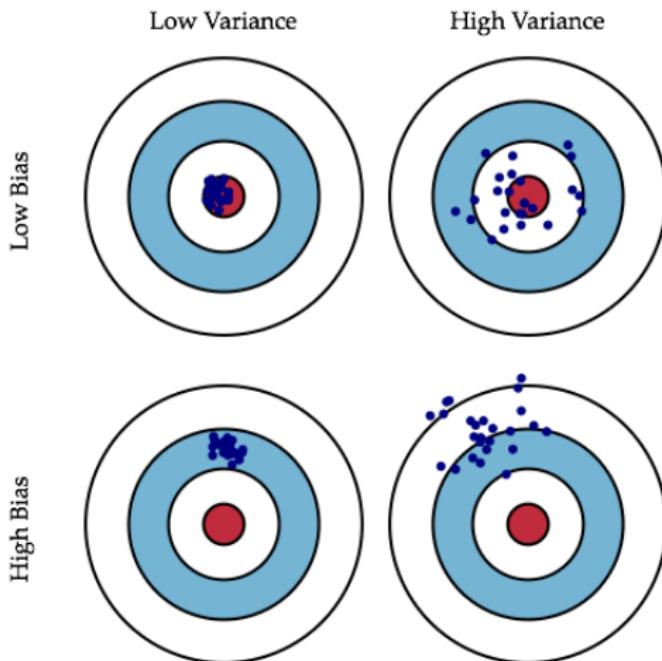


Figure: Graphical illustration of bias vs. variance

# What to do with large bias?

Diagnosis:

- ▶ If your model cannot even fit the training examples, then you have large bias (underfitting)
- ▶ If you can fit the training data, but large error on test data, then you probably have large variance (overfitting)

For bias: redesign your model

- ▶ Add more features as input
- ▶ A more complex model

# What to do with large variance?

- ▶ More data (very effective, but not always practical)
- ▶ Regularization

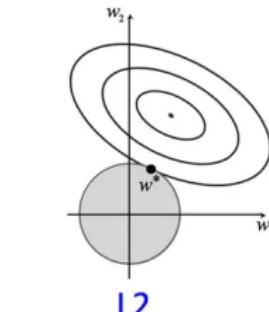
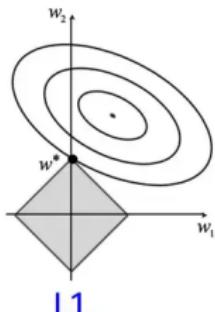
# Regularization

- ▶ L-2 regularization (Ridge regression)

$$L_2(X, y, \mathbf{w}) = \sum_{n=1}^N (f_{\mathbf{w}}(x_n, y_n, \mathbf{w}) - y_n)^2 + \underbrace{\alpha \mathbf{w}^T \mathbf{w}}_{\text{L2 regularization}}$$

- ▶ L-1 regularization (LASSO regression)

$$L_1(X, y, \mathbf{w}) = \sum_{n=1}^N (f_{\mathbf{w}}(x_n, y_n, \mathbf{w}) - y_n)^2 + \underbrace{\alpha |\mathbf{w}|}_{\text{L1 regularization}}$$



# Model Selection

- ▶ Usually a trade-off between bias and variance
- ▶ Select a model that balances two kinds of error to minimize total error
- ▶ Cross-validation: it allows us to estimate the generalization error based on training examples alone.



Figure: Training vs. Validation vs. Test

# K-fold Cross Validation

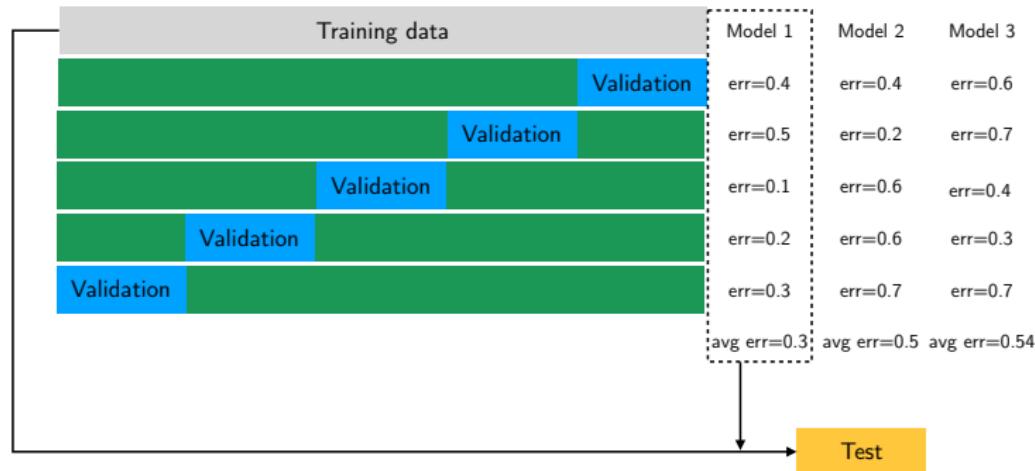


Figure: K-fold Cross Validation

- ▶ In general: we perform K runs, Each run, it allows to use  $(K-1)/K$  of the available data for training.
- ▶ If the number of data is very limited, we can set  $K=N$  (total number of data points). This gives the leave-one-out cross-validation technique.

# The Curse of Dimensionality - when $d$ is large

- ▶ As dimensionality grows: fewer observations per region.
- ▶ When a measure such as a Euclidean distance is defined using many coordinates, there is little difference in the distances between different pairs of samples.
- ▶ The proportion of an inscribed hypersphere with radius  $r$  and dimension  $d$ , to that of a hypercube with edges of length  $2r$ :

$$\frac{V_{\text{hypersphere}}}{V_{\text{hypercube}}} = \frac{\frac{2r^d \pi^{d/2}}{d\Gamma(d/2)}}{(2r)^d} = \frac{\pi^{d/2}}{d2^{d-1}\Gamma(d/2)} \rightarrow 0 \text{ as } d \rightarrow \infty$$

The hypersphere becomes an insignificant volume relative to that of the hypercube

## The Curse of Dimensionality - another angle

Consider a hypersphere of radius  $r = 1$  in a space of  $d$  dimensions, and ask what is the fraction of the volume of the hypersphere that lies between radius  $r = 1 - \epsilon$  and  $r = 1$ . We can evaluate this fraction by noting that the volume of a hypersphere of radius  $r$  in  $d$  dimensions must scale as  $r^d$ , and so we write:

$$V_d(r) = K_d r^d$$

where  $k_d$  depends on  $d$ . Thus the fraction is:

$$\frac{V_d(1) - V_d(1 - \epsilon)}{V_d(1)} = 1 - (1 - \epsilon)^d$$

Most of the volume of a hypersphere is concentrated in a thin shell near the surface!

# The Curse of Dimensionality

- ▶ Real data will often be confined to a region of the space having **lower effective dimensionality**, and in particular the directions over which important variations in the target variables occur may be so confined.
- ▶ Real data will typically exhibit **some smoothness properties** (at least locally) so that for the most part small changes in the input variables will produce small changes in the target variables, and so we can exploit local interpolation-like techniques to allow us to make predictions of the target variables for new values of the input variables.

# Warning of Math

# Maximum Likelihood

We assume there is a true function  $f(x)$  and the true value is given by  $y = f(x) + \epsilon$  where  $\epsilon$  is a Gaussian distribution with mean 0 and variance  $\sigma^2$ . Thus we can write:

$$p(y|x, \mathbf{w}, \beta) = \mathcal{N}(y|f(x), \sigma^2)$$

Assuming the data points are drawn independently from the distribution

- ▶ The likelihood function:  $p(\mathbf{y}|\mathbf{X}, \mathbf{w}, \beta) = \prod_{n=1}^N \mathcal{N}(y_n|f(x_n), \sigma^2)$

# Maximum Likelihood

The likelihood function:

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}, \beta) = \prod_{n=1}^N \mathcal{N}(y_n | f(x_n), \sigma^2)$$

The log-likelihood function:

$$\log p(\mathbf{y}|\mathbf{X}, \mathbf{w}, \beta) = \sum_{n=1}^N \log \mathcal{N}(y_n | f(x_n), \sigma^2) \quad (1)$$

$$= \sum_{n=1}^N \log \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(y_n - f(x_n))^2}{2\sigma^2}} \quad (2)$$

$$= \underbrace{\frac{N}{2} \log \frac{1}{\sigma^2} - \frac{N}{2} \log(2\pi)}_{\text{not related to } \mathbf{w}} - \frac{1}{\sigma^2} E_D(\mathbf{w}) \quad (3)$$

where:

$$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y_n - f(x_n)\}^2$$

# Maximum A Posterior (MAP) Estimation

- ▶ Assuming  $\mathbf{w}$  has a prior distribution
- ▶ We estimate  $p(\mathbf{w}|\mathbf{x}, y)$  which is equivalent to  $\underbrace{p(y|\mathbf{x}, \mathbf{w})}_{\text{likelihood}} \times \underbrace{p(\mathbf{w})}_{\text{prior}}$
- ▶ MAP is equivalent to minimizing the **regularized** sum-of-squares loss under certain assumptions.

# Summary

- ▶ Maximum likelihood (ML)  $\approx$  minimizing the sum-of-squares
- ▶ MAP  $\approx$  minimizing the regularized sum-of-squares
- ▶ Solving for  $\mathbf{w}$  in linear regression when  $f(\mathbf{x}_n) = \mathbf{w}^\top \mathbf{x}_n$ :

$$\mathbf{w}_{ML} = (X^T X)^{-1} X^T \mathbf{y}$$

where  $X$  is an  $N \times d$  matrix. The solution is a linear function of the outputs  $\mathbf{y}$

# End of Warning

## Acknowledgements

Slides adapted from Dr. Percy Liang and Dr. Stefano Ermon's AI at Stanford University and Dr. Bert Huang's Machine Learning at Virginia Tech.