

# CS 559 Machine Learning

## Logistic Regression

Yue Ning  
Department of Computer Science  
Stevens Institute of Technology

# Plan for today

Probabilistic Generative Models

Logistic Regression

# Review of last lecture

- ▶ Bayesian Decision Theory
- ▶ Linear classifiers:
  1. Least square classification,
  2. Fisher's linear discriminant ([Geometrical properties of Linear Discriminant Analysis](#))
  3. Perceptron

# Review of gradient descent algorithms

- ▶ Three important elements: **data( $x, y$ )**, **loss function**, **model parameters  $w$**
- ▶ One important line, gradient update:

$$w^{(t+1)} \leftarrow w^{(t)} - \eta \nabla w$$

# Gradient Descent

---

```
1: procedure BATCH GRADIENT DESCENT
2:   for  $i$  in range(epochs) do
3:      $g^{(i)}(\mathbf{w}) = \text{evaluate\_gradient}(\text{TrainLoss}, \text{data}, \mathbf{w})$ 
4:      $\mathbf{w} = \mathbf{w} - \text{learning\_rate} * g^{(i)}(\mathbf{w})$ 
5:   end for
6: end procedure
```

---

- ▶ Can be very slow.
- ▶ Intractable for datasets that don't fit in memory.
- ▶ Doesn't allow us to update our model online, i.e. with new examples on-the-fly.
- ▶ Guaranteed to converge to the global minimum for convex error surfaces and to a local minimum for non-convex surfaces.

# Stochastic Gradient Descent

---

```
1: procedure STOCHASTIC GRADIENT DESCENT
2:   for  $i$  in range(epochs) do
3:     np.random.shuffle(data)
4:     for example  $\in$  data do
5:        $g^{(i)}(\mathbf{w}) = \text{evaluate\_gradient}(\text{loss}, \text{example}, \mathbf{w})$ 
6:        $\mathbf{w} = \mathbf{w} - \text{learning\_rate} * g^{(i)}(\mathbf{w})$ 
7:     end for
8:   end for
9: end procedure
```

---

- ▶ Allow for online update with new examples.
- ▶ With a high variance that cause the objective function to fluctuate heavily.

# Mini-batch Gradient Descent

---

```
1: procedure MINI-BATCH GRADIENT DESCENT
2:   for  $i$  in range(epochs) do
3:     np.random.shuffle(data)
4:     for batch  $\in$  get_batches(data, batch_size=50) do
5:        $g^{(i)}(\mathbf{w}) = \text{evaluate\_gradient}(\text{loss}, \text{batch}, \mathbf{w})$ 
6:        $\mathbf{w} = \mathbf{w} - \text{learning\_rate} * g^{(i)}(\mathbf{w})$ 
7:     end for
8:   end for
9: end procedure
```

---

- ▶ Reduces the variance of the parameter updates, which can lead to more stable convergence;
- ▶ Can make use of highly optimized matrix optimizations common to state-of-the-art deep learning libraries that make computing the gradient w.r.t. a mini-batch very efficient.

# Evaluation Metrics for Binary Classification

- ▶ The contingency table or confusion matrix:

		True value	
		Positive	Negative
Predicted	Positive	True positive (TP)	False positive (FP)
	Negative	False negative (FN)	True negative (TN)

- ▶ Recall =  $\frac{tp}{tp+fn}$
- ▶ Precision =  $\frac{tp}{tp+fp}$
- ▶ Accuracy =  $\frac{tp+tn}{tp+tn+fp+fn}$  (not a good metric for imbalanced dataset)

# Part I: Probabilistic Generative Models

# Probabilistic Generative Models

- ▶ We now turn to a **probabilistic** approach to classification.
- ▶ How models with linear decision boundaries arise from simple assumptions about the distribution of the data.

# Generative Approach

- ▶ Infer the prior class probabilities  $p(C_k)$ .
- ▶ Estimate the class-conditional densities  $p(\mathbf{x}|C_k)$  for each class  $C_k$ .
- ▶ Use Bayes' theorem to find the class posterior probabilities:

$$p(C_k|\mathbf{x}) = \frac{p(\mathbf{x}|C_k)p(C_k)}{p(\mathbf{x})} \quad (1)$$

where

$$p(\mathbf{x}) = \sum_k p(\mathbf{x}|C_k)p(C_k) \quad (2)$$

- ▶ Use decision theory to determine class membership for each new input  $\mathbf{x}$ .

# Probabilistic Generative Approach-Two Class

Two-class case:

$$\begin{aligned} p(C_1|\mathbf{x}) &= \frac{p(\mathbf{x}|C_1)p(C_1)}{p(\mathbf{x}|C_1)p(C_1) + p(\mathbf{x}|C_2)p(C_2)} \\ &= \frac{1}{1 + e^{-a}} = \sigma(a) \end{aligned} \quad (3)$$

where

$$a = \ln \frac{p(\mathbf{x}|C_1)p(C_1)}{p(\mathbf{x}|C_2)p(C_2)} \quad (4)$$

$$\sigma(a) = \frac{1}{1 + e^{-a}} \text{ (logistic sigmoid function)} \quad (5)$$

## Probabilistic Generative Approach- $K > 2$ Class

$K > 2$  classes:

$$p(C_k|\mathbf{x}) = \frac{p(\mathbf{x}|C_k)p(C_k)}{\sum_{j=1}^K p(\mathbf{x}|C_j)p(C_j)} = \frac{e^{a_k}}{\sum_{j=1}^K e^{a_j}} \quad (6)$$

where

$$a_k = \ln p(\mathbf{x}|C_k)p(C_k) \quad (7)$$

$$\text{softmax}(a_k) = \frac{e^{a_k}}{\sum_{j=1}^K e^{a_j}} \quad (\text{softmax function}) \quad (8)$$

# Probabilistic Generative Models

- ▶ Lets assume the **class-conditional densities** are Gaussian with the same covariance matrix:

$$p(\mathbf{x}|C_k) = \frac{1}{(2\pi)^{D/2}} \frac{1}{|\Sigma|^{1/2}} e^{\left(-\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu}_k)^T \Sigma^{-1} (\mathbf{x}-\boldsymbol{\mu}_k)\right)} \quad (9)$$

- ▶ Two class case first. We can show the following result:

$$P(C_1|\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + \omega_0) \quad (10)$$

$$\mathbf{w} = \Sigma^{-1}(\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2) \quad (11)$$

$$\omega_0 = -\frac{1}{2}\boldsymbol{\mu}_1^T \Sigma^{-1} \boldsymbol{\mu}_1 + \frac{1}{2}\boldsymbol{\mu}_2^T \Sigma^{-1} \boldsymbol{\mu}_2 + \ln \frac{p(C_1)}{p(C_2)} \quad (12)$$

- ▶ How?

# Probabilistic Generative Models

$$P(C_1|\mathbf{x}) = \sigma(a) = \frac{1}{1 + e^{-a}} \quad (13)$$

$$\begin{aligned} a &= \ln \underset{\text{Replace with Eq.9}}{p(\mathbf{x}|C_1)} - \ln \underset{\text{Replace with Eq.9}}{p(\mathbf{x}|C_2)} + \ln \frac{p(C_1)}{p(C_2)} \\ &= -\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu}_1)^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}_1) + \frac{1}{2}(\mathbf{x} - \boldsymbol{\mu}_2)^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}_2) \\ &\quad + \ln \frac{p(C_1)}{p(C_2)} \end{aligned} \quad (14)$$

# Probabilistic Generative Models

Because we want to have  $a = \mathbf{w}^T \mathbf{x} + \omega_0$ , we shift the notations around (see appendix Eq. 45), and get:

$$\mathbf{w} = \Sigma^{-1}(\mu_1 - \mu_2) \quad (15)$$

$$\omega_0 = -\frac{1}{2}\mu_1^T \Sigma^{-1} \mu_1 + \frac{1}{2}\mu_2^T \Sigma^{-1} \mu_2 + \ln \frac{p(C_1)}{p(C_2)} \quad (16)$$

# Probabilistic Generative Models

- We have shown:

$$P(C_1|\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + \omega_0)$$

$$\mathbf{w} = \Sigma^{-1}(\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2)$$

$$\omega_0 = -\frac{1}{2}\boldsymbol{\mu}_1^T \Sigma^{-1} \boldsymbol{\mu}_1 + \frac{1}{2}\boldsymbol{\mu}_2^T \Sigma^{-1} \boldsymbol{\mu}_2 + \ln \frac{p(C_1)}{p(C_2)}$$

- Decision boundary:

$$p(C_1|\mathbf{x}) = p(C_2|\mathbf{x}) = 0.5 \quad (17)$$

$$\Rightarrow \frac{1}{1 + e^{-(\mathbf{w}^T \mathbf{x} + \omega_0)}} = 0.5 \Rightarrow \mathbf{w}^T \mathbf{x} + \omega_0 = 0 \quad (18)$$

# Probabilistic Generative Models

- ▶  $K > 2$  classes:

$$P(C_k|\mathbf{x}) = \frac{p(\mathbf{x}|C_k)p(C_k)}{\sum_j p(\mathbf{x}|C_j)p(C_j)} = \frac{e^{a_k}}{\sum_j e^{a_j}} \quad (19)$$

$$a_k = \ln p(\mathbf{x}|C_k)p(C_k) \quad (20)$$

- ▶ We can show the following result:

$$p(C_k|\mathbf{x}) = \frac{e^{\mathbf{w}_k^T \mathbf{x} + \omega_{k0}}}{\sum_j e^{\mathbf{w}_j^T \mathbf{x} + \omega_{j0}}} \quad (21)$$

$$\mathbf{w}_k = \Sigma^{-1} \boldsymbol{\mu}_k \quad (22)$$

$$\omega_{k0} = -\frac{1}{2} \boldsymbol{\mu}_k^T \Sigma^{-1} \boldsymbol{\mu}_k + \ln p(C_k) \quad (23)$$

# Maximum Likelihood Solution

- We have a parametric functional form for the class-conditional densities:

$$p(\mathbf{x}|C_k) = \frac{1}{(2\pi)^{D/2}} \frac{1}{|\Sigma|^{1/2}} e^{\left(-\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu}_k)^T \boldsymbol{\Sigma}^{-1} (\mathbf{x}-\boldsymbol{\mu}_k)\right)} \quad (24)$$

- We can estimate the parameters and the prior class probabilities using maximum likelihood.

- Two class case with shared covariance matrix.
- Training data:

$$\{\mathbf{x}_n, y_n\}, \quad n = 1, \dots, N$$

$y_n = 1$  denotes class  $C_1$ ;  $y_n = 0$  denotes class  $C_2$ ;

Priors:  $p(C_1) = \gamma, p(C_2) = 1 - \gamma$

# Maximum Likelihood Solution

- ▶ For a data point  $\mathbf{x}_n$  from class  $C_1$ , we have  $y_n = 1$  and therefore:

$$p(\mathbf{x}_n, C_1) = p(\mathbf{x}|C_1)p(C_1) = \gamma \mathcal{N}(\mathbf{x}_n|\boldsymbol{\mu}_1, \Sigma) \quad (25)$$

- ▶ For a data point  $\mathbf{x}_n$  from class  $C_2$ , we have  $y_n = 0$  and therefore:

$$p(\mathbf{x}_n, C_2) = p(\mathbf{x}|C_2)p(C_2) = (1 - \gamma) \mathcal{N}(\mathbf{x}_n|\boldsymbol{\mu}_2, \Sigma) \quad (26)$$

- ▶ Assuming observations are drawn independently, the **likelihood function** is as below:

$$\begin{aligned} p(\mathcal{D}|\gamma, \boldsymbol{\mu}_1, \boldsymbol{\mu}_2, \Sigma) &= \prod_{n=1}^N [p(\mathbf{x}_n, C_1)]^{y_n} [p(\mathbf{x}_n, C_2)]^{1-y_n} \\ &= \prod_{n=1}^N [\gamma \mathcal{N}(\mathbf{x}_n|\boldsymbol{\mu}_1, \Sigma)]^{y_n} [(1 - \gamma) \mathcal{N}(\mathbf{x}_n|\boldsymbol{\mu}_2, \Sigma)]^{1-y_n} \end{aligned}$$

# Maximum Likelihood Solution

- We want to find the values of the parameters that **maximize the likelihood function**, i.e., fit a model that best describes the observed data.

$$p(\mathcal{D}|\gamma, \boldsymbol{\mu}_1, \boldsymbol{\mu}_2, \Sigma) = \prod_{n=1}^N [\gamma \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_1, \Sigma)]^{y_n} [(1 - \gamma) \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_2, \Sigma)]^{1-y_n} \quad (27)$$

- As usual, we consider the log of the likelihood:

$$\begin{aligned} \ln p(\mathcal{D}|\gamma, \boldsymbol{\mu}_1, \boldsymbol{\mu}_2, \Sigma) &= \sum_{n=1}^N [y_n \ln \gamma + y_n \ln \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_1, \Sigma) \\ &\quad + (1 - y_n) \ln (1 - \gamma) + (1 - y_n) \ln \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_2, \Sigma)] \end{aligned} \quad (28)$$

# Maximum Likelihood Solution - parameter $\gamma$

Log Likelihood:

$$\ln p(\mathcal{D}|\gamma, \mu_1, \mu_2, \Sigma) = \sum_{n=1}^N \left[ y_n \ln \gamma + y_n \ln \mathcal{N}(\mathbf{x}_n | \mu_1, \Sigma) + (1 - y_n) \ln (1 - \gamma) + (1 - y_n) \ln \mathcal{N}(\mathbf{x}_n | \mu_2, \Sigma) \right]$$

- We first maximize the log-likelihood with respect to  $\gamma$  (set derivative to 0)

$$\gamma = \frac{1}{N} \sum_{n=1}^N y_n = \frac{N_1}{N} = \frac{N_1}{N_1 + N_2} \quad (29)$$

- The maximum likelihood estimate of  $\gamma$  is the fraction of points in class  $C_1$ . (For multi-class: ML estimate for  $p(C_k)$  is given by the fraction of points in the training set in  $C_k$ ).

# Maximum Likelihood Solution - parameter $\mu$

Log Likelihood:

$$\ln p(\mathcal{D}|\gamma, \boldsymbol{\mu}_1, \boldsymbol{\mu}_2, \Sigma) = \sum_{n=1}^N [y_n \ln \gamma + y_n \ln \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_1, \Sigma) \\ + (1 - y_n) \ln (1 - \gamma) + (1 - y_n) \ln \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_2, \Sigma)]$$

- We then maximize the log-likelihood with respect to  $\boldsymbol{\mu}_1$  (set derivative to 0)

$$\boldsymbol{\mu}_1 = \frac{1}{N_1} \sum_{n=1}^{N_1} y_n \mathbf{x}_n \quad (30)$$

- The maximum likelihood estimate of  $\boldsymbol{\mu}_1$  is the sample mean of all input  $\mathbf{x}_n$  in class  $C_1$ . Same for  $\boldsymbol{\mu}_2$ .

## Maximum Likelihood Solution - parameter $\Sigma$

- ▶ Maximize the log-likelihood with respect to  $\Sigma$  (set derivative to 0), we obtain the estimate  $\Sigma_{ML}$

$$\Sigma_{ML} = \frac{N_1}{N} S_1 + \frac{N_2}{N} S_2 \quad (31)$$

where



$$S_1 = \frac{1}{N_1} \sum_{n \in C_1} (\mathbf{x}_n - \boldsymbol{\mu}_1)(\mathbf{x}_n - \boldsymbol{\mu}_1)^T \quad (32)$$

$$S_2 = \frac{1}{N_2} \sum_{n \in C_2} (\mathbf{x}_n - \boldsymbol{\mu}_2)(\mathbf{x}_n - \boldsymbol{\mu}_2)^T \quad (33)$$

- ▶ The maximum likelihood estimate of the covariance is given by the weighted average of the sample covariance matrices associated with each of the classes.
- ▶ The results extend to K classes.

# Summary so far

- ▶ We assumed  $p(\mathbf{x}|(y=1)) \sim \mathcal{N}(\boldsymbol{\mu}_+, \boldsymbol{\Sigma}_+)$  and  $p(\mathbf{x}|(y=-1)) \sim \mathcal{N}(\boldsymbol{\mu}_-, \boldsymbol{\Sigma}_-)$ , and two class-probabilities  $p(y=1)$  and  $p(y=-1)$ .
- ▶ This is called an **generative model**, as we have written down a full joint model over the data.
- ▶ We saw that violations of the model assumption can lead to “bad” decision boundaries.



Figures from Bishop PRML, 44a and b

# Probabilistic Generative Models - parameters

- ▶ How many parameters did we estimate to fit Gaussian class-conditional densities (the generative approach)?

$$p(C_1) \Rightarrow 1$$

2 mean vectors  $\Rightarrow 2d$

$$\Sigma \Rightarrow d + \frac{d^2 - d}{2} = \frac{d^2 + d}{2}$$

$$\text{total} = 1 + 2d + \frac{d^2 + d}{2} = O(d^2)$$

## Part II: Logistic Regression

# Logistic Regression

$$P(C_1|\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + \omega_0) = f(\mathbf{x}) \quad (34)$$

- We use maximum likelihood to determine the parameters of the logistic regression model.

$$\{\mathbf{x}_n, y_n\}, \quad n = 1, \dots, N$$

$y_n = 1$  denotes class  $C_1$ ;  $y_n = 0$  denotes class  $C_2$ ;

- We want to find the values of  $\mathbf{w}$  that maximize the posterior probabilities associated to the observed data
- Likelihood function:

$$\begin{aligned} \mathcal{L}(\mathbf{w}) &= \prod_{n=1}^N p(C_1|\mathbf{x}_n)^{y_n} (1 - p(C_1|\mathbf{x}_n))^{1-y_n} \\ &= \prod_{n=1}^N f(\mathbf{x}_n)^{y_n} (1 - f(\mathbf{x}_n))^{1-y_n} \end{aligned} \quad (35)$$

# Logistic Regression

$$P(C_1|\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + \omega_0) = f(\mathbf{x})$$


- We consider the negative logarithm of the likelihood (Cross Entropy):

$$\begin{aligned}\mathcal{E}(\mathbf{w}) &= -\ln \mathcal{L}(\mathbf{w}) = -\ln \prod_{n=1}^N p(C_1|\mathbf{x}_n)^{y_n} (1 - p(C_1|\mathbf{x}_n))^{1-y_n} \\ &= - \sum_{n=1}^N (y_n \ln f(\mathbf{x}_n) + (1 - y_n) \ln (1 - f(\mathbf{x}_n)))\end{aligned}\tag{36}$$

- Thus:

$$\max \mathcal{L}(\mathbf{w}) = \min \mathcal{E}(\mathbf{w}) \tag{37}$$

The cost-function for logistic regression is convex.



- ▶ Fact: The negative log-likelihood is *convex* – this makes life much more easier.
- ▶ There are no local minima to get stuck in, and there is good optimization techniques for convex problems.

# Logistic Regression - Gradient

$$P(C_1|\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + \omega_0) = f(\mathbf{x})$$

- We compute the derivative of the error function with respect to  $\mathbf{w}$

$$\frac{\partial \mathcal{E}(\mathbf{w})}{\partial \mathbf{w}} = \frac{\partial}{\partial \mathbf{w}} \left[ -\ln \prod_{n=1}^N p(C_1|\mathbf{x}_n)^{y_n} (1 - p(C_1|\mathbf{x}_n))^{1-y_n} \right] \quad (38)$$

- The derivative of the logistic sigmoid function:

$$\begin{aligned} \frac{\partial}{\partial a} \sigma(a) &= \frac{\partial}{\partial a} \frac{1}{1 + e^{-a}} = \frac{e^{-a}}{(1 + e^{-a})^2} \\ &= \frac{1}{1 + e^{-a}} \frac{e^{-a}}{(1 + e^{-a})} = \frac{1}{1 + e^{-a}} \left(1 - \frac{1}{1 + e^{-a}}\right) \\ &= \sigma(a)(1 - \sigma(a)) \end{aligned} \quad (39)$$

# Logistic Regression

The derivative of the loss function with respect to  $w$  is ?

# Probabilistic Discriminative Models - parameters

- ▶ Two-class case:

$$P(C_1|\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + \omega_0) = f(\mathbf{x})$$

$$P(C_2|\mathbf{x}) = 1 - P(C_1|\mathbf{x})$$

- ▶ This model is known as Logistic Regression.
- ▶ Assuming  $\mathbf{x} \in \mathbb{R}^d$ , how many parameters do we need to estimate?

$$d + 1$$

## Summary so far

- ▶ From the previous introduction, we know that

$$P(y = 1|\mathbf{x}) = \sigma(a(\mathbf{x}))$$


where  $\sigma(a) = 1/(1 + \exp(-a))$  and  $a(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + \omega_o$ .

- ▶ Notation is simpler if we use 0 and 1 as class labels, so we define  $y_n = 1$  as the label for the positive class, and  $y_n = 0$  as label for the negative class.
- ▶ In other words,  $y|x \sim \text{Bernoulli}(\sigma(f(x)))$ .
- ▶ The parameters of  $a(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + \omega_o$  can be learned by minimizing the negative log-likelihood  
$$L(\mathbf{w}) = -\sum_n \{y_n \log p(y_n|\mathbf{x}_n, \mathbf{w}) + (1-y_n) \log(1-p(y_n|\mathbf{x}_n, \mathbf{w}))\}$$
- ▶ This is a **discriminative** approach to classification, as we only model the labels, and not the inputs.

## Summary: Discriminative vs. Generative

Decision rule and function shape of  $p(y|x)$  will be the same for the generative and the discriminative model, but the parameters were obtained differently.

# Maximum likelihood estimation of Logistic Regression



Bishop PRML Figure 44 a and b

- ▶ Logistic regression is a *much* better algorithm than the algorithms we discussed last week.
- ▶ Need to optimize log-likelihood numerically.
- ▶ People typically minimize the negative log-likelihood  $\mathcal{L}$  rather than maximize the log-likelihood...
- ▶ To numerically minimize the negative log-likelihood, we need its gradient (and maybe its hessian)

# Multiclass Logistic Regression

- Multiclass case:

$$p(C_k | \mathbf{x}) = \frac{e^{\mathbf{w}_k^T \mathbf{x} + \omega_{k0}}}{\sum_j e^{\mathbf{w}_j^T \mathbf{x} + \omega_{j0}}} = f_k(\mathbf{x}) \quad (40)$$

- We use maximum likelihood to determine the parameters:

$$\{\mathbf{x}_n, y_n\}, \quad n = 1, \dots, N$$

$y_n = (0, \dots, 1, \dots, 0)$  denotes class  $C_k$

- We want to find the values of  $\mathbf{w}_1, \dots, \mathbf{w}_k$  that maximize the posterior probabilities associated to the observed data likelihood function:

$$\mathcal{L}(\mathbf{w}_1, \dots, \mathbf{w}_k) = \prod_{n=1}^N \prod_{k=1}^K p(C_k | \mathbf{x}_n)^{y_{nk}} = \prod_{n=1}^N \prod_{k=1}^K f_k(\mathbf{x}_n)^{y_{nk}} \quad (41)$$

# Multiclass Logistic Regression

$$\mathcal{L} = \prod_{n=1}^N \prod_{k=1}^K p(C_k | \mathbf{x}_n)^{y_{nk}} = \prod_{n=1}^N \prod_{k=1}^K f_k(\mathbf{x}_n)^{y_{nk}}$$

- ▶ Consider the negative logarithm of the likelihood:

$$\mathcal{E}(\mathbf{w}_1, \dots, \mathbf{w}_k) = -\ln \mathcal{L}(\mathbf{w}_1, \dots, \mathbf{w}_k) = -\sum_{n=1}^N \sum_{k=1}^K y_{nk} \ln f_k(\mathbf{x}_n) \quad (42)$$

$$\min_{\mathbf{w}_j} \mathcal{E}((\mathbf{w}_j)) \quad (43)$$

- ▶ The gradient of the error function w.r.t one of the parameter vectors:

$$\frac{\partial}{\partial \mathbf{w}_j} \mathcal{E}(\mathbf{w}_1, \dots, \mathbf{w}_k) = \frac{\partial}{\partial \mathbf{w}_j} \left[ -\sum_{n=1}^N \sum_{k=1}^K y_{nk} \ln f_k(\mathbf{x}_n) \right] \quad (44)$$

# Multiclass Logistic Regression

$$\frac{\partial}{\partial \mathbf{w}_j} \mathcal{E}(\mathbf{w}_1, \dots, \mathbf{w}_k) = \frac{\partial}{\partial \mathbf{w}_j} \left[ - \sum_{n=1}^N \sum_{k=1}^K y_{nk} \ln f_k(\mathbf{x}_n) \right]$$

- ▶ The derivatives of the softmax function:

$$\frac{\partial}{\partial a_k} f_k = \frac{\partial}{\partial a_k} \frac{e^{a_k}}{\sum_j e^{a_j}} = \frac{e^{a_k} \sum_j e^{a_j} - e^{a_k} e^{a_k}}{(\sum_j e^{a_j})^2} = f_k - f_k^2 = f_k(1-f_k)$$

- ▶ Thus:

$$\text{for } j \neq k, \frac{\partial}{\partial a_j} f_k = \frac{\partial}{\partial a_j} f_k = \frac{\partial}{\partial a_j} \frac{e^{a_k}}{\sum_j e^{a_j}} = \frac{-e^{a_k} e^{a_j}}{(\sum_j e^{a_j})^2} = -f_k f_j$$

- ▶ Compact expression ( $I_{kj}$  are the elements of the identity matrix)

$$\frac{\partial}{\partial a_j} f_k = f_k(I_{kj} - f_j)$$

# Multiclass Logistic Regression

$$\nabla_{\mathbf{w}_j} \mathcal{E}(\mathbf{w}_1, \dots, \mathbf{w}_k) = \sum_{n=1}^N (f_{nj} - y_{nj}) \mathbf{x}_n$$

# Multiclass Logistic Regression

$$\nabla_{\mathbf{w}_j} \mathcal{E}(\mathbf{w}_1, \dots, \mathbf{w}_k) = \sum_{n=1}^N (f_{nj} - y_{nj}) \mathbf{x}_n$$

- ▶ It can be shown that  $\mathcal{E}$  is a convex function of  $\mathbf{w}$ . Thus, it has a unique minimum.
- ▶ For a batch solution, we can use the Newton-Raphson optimization technique.
- ▶ Online solution (SGD):

$$\mathbf{w}_j^{t+1} = \mathbf{w}_j^t - \eta \nabla_{\mathbf{w}_j} \mathcal{E}_n(\mathbf{w}) = \mathbf{w}_j^t - \eta (f_{nj} - y_{nj}) \mathbf{x}_n$$

## Acknowledgements

Slides adapted from Dr. Carlotta Domeniconi's *Pattern Recognition* at George Mason University.

# Appendix

$$\begin{aligned} a &= -\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu}_1)^T \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu}_1) + \frac{1}{2}(\mathbf{x} - \boldsymbol{\mu}_2)^T \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu}_2) + \ln \frac{p(C_1)}{p(C_2)} \\ &= -\frac{1}{2}\mathbf{x}^T \Sigma^{-1} \mathbf{x} + \frac{1}{2}\boldsymbol{\mu}_1^T \Sigma^{-1} \mathbf{x} + \frac{1}{2}\mathbf{x}^T \Sigma^{-1} \boldsymbol{\mu}_1 - \frac{1}{2}\boldsymbol{\mu}_1^T \Sigma^{-1} \boldsymbol{\mu}_1 \\ &\quad + \frac{1}{2}\mathbf{x}^T \Sigma^{-1} \mathbf{x} - \frac{1}{2}\boldsymbol{\mu}_2^T \Sigma^{-1} \mathbf{x} - \frac{1}{2}\mathbf{x}^T \Sigma^{-1} \boldsymbol{\mu}_2 + \frac{1}{2}\boldsymbol{\mu}_2^T \Sigma^{-1} \boldsymbol{\mu}_2 + \ln \frac{p(C_1)}{p(C_2)} \\ &= \boldsymbol{\mu}_1^T \Sigma^{-1} \mathbf{x} - \boldsymbol{\mu}_2^T \Sigma^{-1} \mathbf{x} - \frac{1}{2}\boldsymbol{\mu}_1^T \Sigma^{-1} \boldsymbol{\mu}_1 + \frac{1}{2}\boldsymbol{\mu}_2^T \Sigma^{-1} \boldsymbol{\mu}_2 + \ln \frac{p(C_1)}{p(C_2)} \\ &= (\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2)^T \Sigma^{-1} \mathbf{x} + \omega_0 = \mathbf{w}^T \mathbf{x} + \omega_0 \end{aligned} \tag{45}$$