

# CS 559 Machine Learning

## Support Vector Machines

Yue Ning  
Department of Computer Science  
Stevens Institute of Technology

# Plan for today


Vector algebra, Formulation, Margin

Non-separable Case, Penalties

Non-linearity, Kernels

# Linear Classifier

- ▶ Linear classifiers construct linear decision boundaries(hyperplanes) that try to separate the data into different classes as well as possible.



- ▶ Classification Rule (the Perceptron model):

$$\text{Input: } \mathbf{x} \in \mathbb{R}^d \quad \text{Output: } \text{sign}(\mathbf{w}^\top \mathbf{x} + b)$$

- ▶ The classifier computes a linear combination of the input features, and return the sign.

# Perceptron Learning Algorithm as Gradient Descent

Objective: Find a separating hyperplane that correctly classifier all input patterns.

- ▶ There are two types of error:

$$y_i = 1 \text{ and } \mathbf{w}^\top \mathbf{x}_i + b < 0$$

$$y_i = -1 \text{ and } \mathbf{w}^\top \mathbf{x}_i + b > 0$$

- ▶ Thus, for all misclassified points:

$$y_i(\mathbf{w}^\top \mathbf{x}_i + b) < 0$$

- ▶ To reduce the number of misclassified points, the goal is to minimize:

$$D(\mathbf{w}, b) = - \sum_{i \in M} y_i(\mathbf{w}^\top \mathbf{x}_i + b)$$

where M indexes the set of misclassified points.

# Perceptron Learning Algorithm as Gradient Descent

Objective: Find a separating hyperplane that correctly classifies all input patterns.

- ▶ To minimize  $D(\mathbf{w}, b)$  we can perform gradient descent over the surface represented by  $D(\mathbf{w}, b)$  in parameter space. We iteratively move along the opposite direction of the gradient till a minimum is reached.
- ▶ The gradient is given by:

$$\frac{\partial D(\mathbf{w}, b)}{\partial \mathbf{w}} = - \sum_{i \in M} y_i \mathbf{x}_i$$

$$\frac{\partial D(\mathbf{w}, b)}{\partial b} = - \sum_{i \in M} y_i$$

- ▶ Update rule for parameters (“steepest descent”):

$$\mathbf{w}' = \mathbf{w} + \sum_{i \in M} y_i \mathbf{x}_i$$

$$b' = b + \sum_{i \in M} y_i$$

# Perceptron Learning Algorithm as Gradient Descent

- ▶ In practice, “Stochastic Gradient Descent (SGD)” is used: rather than computing the sum of the gradient contributions of each  $x_i$ , followed by a step in the negative gradient direction, a step is taken after each observation is visited. The resulting update rule for parameters is:

$$\mathbf{w}' = \mathbf{w} + y_i \mathbf{x}_i$$


$$b' = b + y_i$$

# Problems with Perceptron Learning Algorithm

- ▶ When the data are **linearly separable**, there are many solutions, and which one is found depends on the starting values of the parameters.
- ▶ The finite **number of steps** to convergence can be very large: the smaller the gap between the two classes, the longer the time to find it.
- ▶ When the data are **not separable**, the algorithm will **not converge**, and cycles develop. The cycles can be long, and therefore hard to detect. We can provide a solution to the first problem by adding additional constraints to the separating hyperplane we want to find!

# Some Vector Algebra

- ▶ Hyperplane  $L$ :  $f(\mathbf{x}) = (\mathbf{w}^\top \mathbf{x} + b) = 0$ , when  $\mathbf{x} \in \mathbb{R}^2$ ,  $f(\mathbf{x})$  is a line.



## Some Vector Algebra - property 1

- ▶ Consider any two points  $x_1, x_2$ , lying on hyperplane L:

$$\begin{aligned}\mathbf{w}x_1 + b &= 0 \\ \mathbf{w}x_2 + b &= 0\end{aligned}\rightarrow \mathbf{w}^\top(x_1 - x_2) = 0$$

- ▶ Since  $\mathbf{w}^\top(x_1 - x_2) = \mathbf{w} \cdot (x_1 - x_2) = 0$ , the two vectors  $\mathbf{w}$  and  $x_1 - x_2$  are orthogonal vectors.

$$\mathbf{w}^* = \frac{\mathbf{w}}{\|\mathbf{w}\|}$$

is the vector normal to the surface of L.

- ▶ Note 1: All vectors here are column vectors.
- ▶ Note 2: Dot product (inner product) of two vectors  $\mathbf{a} \cdot \mathbf{b} = \mathbf{a}^\top \mathbf{b} = \|\mathbf{a}\| \times \|\mathbf{b}\| \times \cos \alpha$  where  $\alpha$  is the angle between  $a$  and  $b$ .

## Some Vector Algebra - property 2


- ▶ For any point  $\mathbf{x}_0$  on  $L$ :

$$\mathbf{w}^\top \mathbf{x}_0 + b = 0$$

Thus:

$$\mathbf{w}^\top \mathbf{x}_0 = -b$$

## Some Vector Algebra - property 3




The signed distance of any point  $\underline{x}$  to  $L$  is:

$$(\mathbf{w}^*)^\top (\underline{x} - \underline{x}_0) = \frac{\mathbf{w}^\top}{\|\mathbf{w}\|} (\underline{x} - \underline{x}_0) = \frac{1}{\|\mathbf{w}\|} (\mathbf{w}^\top \underline{x} - \mathbf{w}^\top \underline{x}_0) = \frac{1}{\|\mathbf{w}\|} (\mathbf{w}^\top \underline{x} + b)$$

# Largest Margin Hyperplanes

Goal: Find the hyperplane that separates the two classes and maximizes the distance to the closest points from either class.

- ▶ Such distance is called margin.



- ▶ Provide a unique solution to the separating hyperplane problem;
- ▶ Maximizing the margin between the two classes on the training data gives better classification performance on test data.

# The Training Data

For two classes:

$$(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)$$

$$\mathbf{x}_i \in \mathbf{R}^d$$

$$y_i = \{-1, +1\}$$

We need to formalize the largest margin criterion.

# Formulation

Consider the following optimization problem:

$$\max_{\mathbf{w}, b} 2C$$

$$\text{subject to } \frac{1}{\|\mathbf{w}\|} y_i (\mathbf{w}^\top \mathbf{x}_i + b) \geq C \quad i = 1, \dots, N$$

- ▶ Remember Property 3:

The signed distance of any point  $\mathbf{x}$  to L is:

$$\frac{1}{\|\mathbf{w}\|} (\mathbf{w}^\top \mathbf{x}_i + b)$$

- ▶ Thus, the set of conditions above ensure that all the training data are at least at distance  $C$  from the decision boundary.
- ▶ We seek the largest such  $C$  and associated parameters.
- ▶ We can rewrite the above conditions as:

$$y_i (\mathbf{w}^\top \mathbf{x}_i + b) \geq C \|\mathbf{w}\|$$

# Formulation

- We can rewrite the above conditions as:

$$y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq C\|\mathbf{w}\|$$

- Since  $\mathbf{w}^\top \mathbf{x} + b = 0$  and  $c(\mathbf{w}^\top \mathbf{x} + b) = 0$  define the same plane, we can arbitrarily normalize  $\|\mathbf{w}\| = \frac{1}{C}$
- The original maximization problem is equivalent to:

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2$$

subject to  $y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 \quad i = 1, \dots, N$

- The constraints define an empty margin around the linear decision boundary of thickness  $\frac{2}{\|\mathbf{w}\|}$ . We choose  $\mathbf{w}, b$  to maximize its thickness.
- This is a quadratic (convex) optimization problem subject to linear constraints and there is a unique minimum

# Lagrange Multipliers

- ▶ We introduce Lagrange multipliers, which allow to take the constraints within the function to be minimized. Two reasons for doing this:
  1. The constraints will be **replaced** by constraints on the Lagrange multipliers themselves, which are easier to handle.
  2. Training data will only appear **in the form of dot products between vectors** (crucial for nonlinear case).
- ▶ We introduce the **Lagrange multipliers**  $\alpha_i \geq 0$ ,  $i = 1, \dots, N$ , one for each of the inequality constraints.
- ▶ Recall the rule: for constraints of the form  $F_i \geq 0$ , the constraint equations are **multiplied by** Lagrange multipliers and **subtracted from** the objective function, to form the Lagrangian.

# Lagrange Multipliers - Primal Form

- We then obtain the Lagrangian: (also called “primal form”):

$$L_p = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{i=1}^N \alpha_i [y_i (\mathbf{w}^\top \mathbf{x}_i + b) - 1]$$

- We must now minimize  $L_p$  with respect to  $\mathbf{w}$  and  $b$ :

$$\min_{\mathbf{w}, b} \max_{\alpha_i \geq 0} L_p$$

- This indicates that this is the primal form of the optimization problem.
- We will actually solve the optimization problem by solving for the dual of the original problem

## Dual Form

- ▶ The solution to the dual form provides a **lower bound** to the solution of the primal form
- ▶ What is the dual form?

$$\max_{\alpha_i \geq 0} \min_{\mathbf{w}, b} L_p$$

- ▶ Setting the derivatives to zero gives:

$$\frac{\partial L_p}{\partial \mathbf{w}} = \mathbf{w} - \sum_{i=1}^N \alpha_i y_i \mathbf{x}_i = 0 \Rightarrow \mathbf{w} = \sum_{i=1}^N \alpha_i y_i \mathbf{x}_i \quad (1)$$

$$\frac{\partial L_p}{\partial b} = - \sum_{i=1}^N \alpha_i y_i = 0 \Rightarrow \sum_{i=1}^N \alpha_i y_i = 0 \quad (2)$$

# Dual Form

Substituting eq. 1 and 2 in  $L_p$  gives:

$$\begin{aligned} L_D &= \frac{1}{2} \left( \sum_{i=1}^N \alpha_i y_i \mathbf{x}_i \right) \left( \sum_{k=1}^N \alpha_k y_k \mathbf{x}_k \right) - \sum_{i=1}^N \alpha_i \left[ y_i \left( \mathbf{x}_i^\top \left( \sum_{k=1}^N \alpha_k y_k \mathbf{x}_k \right) + b \right) - 1 \right] \\ &= \frac{1}{2} \sum_{i=1}^N \sum_{k=1}^N \alpha_i \alpha_k y_i y_k \mathbf{x}_i^\top \mathbf{x}_k - \sum_{i=1}^N \sum_{k=1}^N \alpha_i \alpha_k y_i y_k \mathbf{x}_i^\top \mathbf{x}_k - b \underbrace{\sum_{i=1}^N \alpha_i y_i}_{=0} + \sum_{i=1}^N \alpha_i \\ &= \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{k=1}^N \alpha_i \alpha_k y_i y_k \mathbf{x}_i^\top \mathbf{x}_k \\ \text{subject to } \alpha_i &\geq 0 \end{aligned}$$

# The Lagrangian Dual Form

- ▶ Solution is obtained by **maximizing**  $L_D$  with respect to the  $\alpha_i$ .
- ▶ Solution must satisfy the conditions:

$$\mathbf{w} = \sum_{i=1}^N \alpha_i y_i \mathbf{x}_i$$


$$\sum_{i=1}^N \alpha_i y_i = 0$$

$$\alpha_i \geq 0$$

$$\alpha_i [y_i (\mathbf{w}^\top \mathbf{x}_i + b) - 1] = 0 \quad \forall i = 1, \dots, N$$

## Dual Form

- ▶ If  $\alpha_i > 0$ , then  $y_i(\mathbf{w}^\top \mathbf{x}_i + b) = 1$ , that is  $\mathbf{x}_i$  is on the boundary of the margin.
- ▶ If  $y_i(\mathbf{w}^\top \mathbf{x}_i + b) > 1$ ,  $\mathbf{x}_i$  is not on the boundary of the margin, and  $\alpha_i = 0$



## Dual Form

- ▶ The solution vector  $\mathbf{w}$  is:  $\mathbf{w} = \sum_{i=1}^N \alpha_i y_i \mathbf{x}_i$ .
- ▶  $\mathbf{x}_i$  are SUPPORT VECTORS when  $\alpha_i > 0$ . We have three support vectors in the above example.
- ▶ To obtain the value of  $b$ : solve  $\alpha_i [y_i (\mathbf{w}^\top \mathbf{x}_i + b) - 1] = 0$  for any of the support vectors.
- ▶ The largest margin hyperplane gives a function:  
 $f(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + b$  for classifying new observations  
 $\hat{y} = \text{sign}(f(\mathbf{x}))$

# Observations

- ▶ The **support vectors** are the critical elements of the training set. They lie closest to the decision boundary.
- ▶ Only the support vectors affect the solution
- ▶ However, the identification of the support vectors **requires the use of all the training data**.

## Summary so far

- ▶ Training data:  $(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)$ ,  
 $\mathbf{x}_i \in \mathbb{R}^d, y_i \in \{-1, +1\}$
- ▶ When the two classes are linearly separable, we can find a function  $f(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + b$  with  $y_i f(\mathbf{x}_i) > 0 \forall i$
- ▶ In particular, we can find the hyperplane that creates the largest margin between the training points.
- ▶ The optimization problem captures this concept

$$\max_{\mathbf{w}, b} 2C$$


$$\text{subject to } \frac{1}{\|\mathbf{w}\|} y_i (\mathbf{w}^\top \mathbf{x}_i + b) \geq C \quad i = 1, \dots, N$$

- ▶ It can be more conveniently rewritten as below where  $C = \frac{1}{\|\mathbf{w}\|}$

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2$$

$$\text{subject to } y_i (\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 \quad i = 1, \dots, N$$

# Geometric perspective



## The Non-separable Case

- ▶ Suppose now the classes overlap. We can still maximize  $C$ , but allow for some points to be on the wrong side of the margin.
- ▶ We need to modify the constraints we had for the separable case:  $\frac{1}{\|\mathbf{w}\|}y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq C \quad i = 1, \dots, N$ .
- ▶ To achieve this goal, we define  $N$  slack variables:

$$\xi_1, \xi_2, \dots, \xi_N$$

- ▶ Then a natural way to modify the constraints above is:

$$\frac{1}{\|\mathbf{w}\|}y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq C(1 - \xi_i) \quad i = 1, \dots, N$$


$$\text{with } \xi_i \geq 0 \quad \forall i, \quad \sum_{i=1}^N \xi_i \leq \text{constant}$$

## The Non-separable Case

- ▶ Idea of the formulation:  $\xi_i$  is the proportional amount by which the prediction  $f(\mathbf{x}_i)$  is on the wrong side of the margin.
- ▶ Note:  $C(1 - \xi_i) = C - C\xi_i$

# Slack Variables

- ▶ A geometric perspective:



- ▶ The points  $(\mathbf{x}_1^*, \mathbf{x}_2^*, \mathbf{x}_3^*, \mathbf{x}_4^*, \mathbf{x}_5^*)$  are on the wrong side of their margin.
- ▶ Point  $\mathbf{x}_i^*$  is on the wrong side of its margin by an amount  $C\xi_i$
- ▶ Point  $\mathbf{x}_j^*$  on the correct side have  $\xi_j = 0$
- ▶ Misclassification occurs when  $\xi_i > 1 \Rightarrow C(1 - \xi_i) < 0$ , e.g., points  $\mathbf{x}_3^*$  and  $\mathbf{x}_5^*$  are misclassified by the given boundary.

## Slack Variables (2)

- ▶ The condition  $\sum_{i=1}^N \xi_i \leq \text{constant}$  bounds the sum  $\sum_i \xi_i$
- ▶ Thus, it bounds the total proportional amount by which predictions fall on the wrong side of their margin.
- ▶ Since misclassification occur when  $\xi_i > 1$  (in this case  $y_i f(\mathbf{x}_i) < 0$ ), bounding  $\sum_{i=1}^N \xi_i \leq k$ , bounds the total number of training misclassifications at  $k$ .
- ▶ So, for the non-separable case, we have the optimization problem:

$$\max_{\mathbf{w}, b} 2C$$

$$\text{subject to: } \frac{1}{\|\mathbf{w}\|} y_i (\mathbf{w}^\top \mathbf{x}_i + b) \geq C(1 - \xi_i) \quad i = 1, \dots, N$$

$$\text{with } \xi_i \geq 0 \quad \forall i, \quad \sum_{i=1}^N \xi_i \leq \text{constant}$$

## Slack Variables (3)

- As for the separable case, we define  $C = \frac{1}{\|\mathbf{w}\|}$  and rewrite the above maximization problem in the equivalent form:

$$\min_{\mathbf{w}, b} \frac{\|\mathbf{w}\|^2}{2}$$

subject to:  $y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 - \xi_i \quad i = 1, \dots, N$

with  $\xi_i \geq 0 \quad \forall i, \quad \sum_{i=1}^N \xi_i \leq \text{constant}$

- We have obtained a quadratic optimization problem with linear constraints. We will solve it using Lagrange multipliers.

# Lagrange Multipliers for Slack Variables

- ▶ First, one more step: we have seen that the condition  $\sum_i \xi_i \leq \text{constant}$ , bounds the number of training misclassifications.
- ▶ We can incorporate this condition into the objective function by adding an extra cost for errors:

$$\min_{\mathbf{w}, b} \frac{\|\mathbf{w}\|^2}{2} + \gamma \sum_{i=1}^N \xi_i$$

subject to:  $y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 - \xi_i \quad i = 1, \dots, N$

with  $\xi_i \geq 0 \quad \forall i, \quad \sum_{i=1}^N \xi_i \leq \text{constant}$

here,  $\gamma$  is a parameter to be chosen by the user. A larger  $\gamma$  corresponds to assigning a higher penalty to errors.

# Lagrange Multipliers for Slack Variables

$$\min_{\mathbf{w}, b} \frac{\|\mathbf{w}\|^2}{2} + \gamma \sum_{i=1}^N \xi_i$$

subject to:  $y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 - \xi_i \quad i = 1, \dots, N$

with  $\xi_i \geq 0 \quad \forall i, \quad \sum_{i=1}^N \xi_i \leq \text{constant}$

- ▶ Introducing the Lagrange multipliers  $\alpha_i$  and  $\mu_i$  (one for each constraint), gives the following Lagrange (primal) function:

$$L_p = \frac{1}{2} \|\mathbf{w}\|^2 + \gamma \sum_{i=1}^N \xi_i - \sum_{i=1}^N \alpha_i [y_i(\mathbf{w}^\top \mathbf{x}_i + b) - (1 - \xi_i)] - \sum_{i=1}^N \mu_i \xi_i$$

- ▶ Our objective is:  $\min_{\mathbf{w}, b, \xi_i} L_p$

# Lagrange Multipliers Solution

- ▶ Setting the respective derivatives to zero gives:

$$\frac{\partial L_p}{\partial \mathbf{w}} = \mathbf{w} - \sum_{i=1}^N \alpha_i y_i \mathbf{x}_i = 0 \Rightarrow \mathbf{w} = \sum_{i=1}^N \alpha_i y_i \mathbf{x}_i \quad (3)$$

$$\frac{\partial L_p}{\partial b} = - \sum_{i=1}^N \alpha_i y_i = 0 \Rightarrow \sum_{i=1}^N \alpha_i y_i = 0 \quad (4)$$

$$\frac{\partial L_p}{\partial \xi_i} = \gamma - \alpha_i - \mu_i \quad \forall i \Rightarrow \alpha_i = \gamma - \mu_i \quad \forall i \quad (5)$$

along with the positivity constraints  $\alpha_i, \mu_i, \xi_i \geq 0 \quad \forall i$

- ▶ Substituting eq. 3, 4, 5 in  $L_p$ , we obtain the so called dual objective function:

$$L_D = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j \mathbf{x}_i^\top \mathbf{x}_j$$

where  $L_D$  gives a lower bound on the objective function

# Deriving the Dual Form

$$\begin{aligned} & \frac{1}{2} \left( \sum_{i=1}^N \alpha_i y_i \mathbf{x}_i \right)^\top \left( \sum_{j=1}^N \alpha_j y_j \mathbf{x}_j \right) + \gamma \sum_{i=1}^N \xi_i - \sum_{i=1}^N \alpha_i y_i \mathbf{x}_i^\top \mathbf{w}^\top - \sum_{i=1}^N \alpha_i y_i b + \sum_i \alpha_i (1 - \xi_i) - \sum_{i=1}^N \mu_i \xi_i \\ &= \frac{1}{2} \sum_i \sum_j \alpha_i \alpha_j y_i y_j \mathbf{x}_i^\top \mathbf{x}_j - \sum_i \sum_j \alpha_i \alpha_j y_i y_j \mathbf{x}_i^\top \mathbf{x}_j + \gamma \sum_{i=1}^N \xi_i - b \underbrace{\sum_{i=1}^N \alpha_i y_i}_{=0} + \sum_i \alpha_i (1 - \xi_i) - \sum_{i=1}^N \mu_i \xi_i \\ &= -\frac{1}{2} \sum_i \sum_j \alpha_i \alpha_j y_i y_j \mathbf{x}_i^\top \mathbf{x}_j + \gamma \sum_{i=1}^N \xi_i + \sum_{i=1}^N \alpha_i - \sum_i \alpha_i \xi_i - \sum_i \mu_i \xi_i \\ &= \sum_i^N \alpha_i - \frac{1}{2} \sum_i \sum_j \alpha_i \alpha_j y_i y_j \mathbf{x}_i^\top \mathbf{x}_j + \underbrace{\sum_i^N (\gamma - \mu_i) \xi_i}_{=0} - \sum_i \alpha_i \xi_i \\ &= \sum_i \alpha_i - \frac{1}{2} \sum_i \sum_j \alpha_i \alpha_j y_i y_j \mathbf{x}_i^\top \mathbf{x}_j \end{aligned}$$

# Lagrange Multipliers Solution

- ▶ Thus: the solution is obtained by maximizing  $L_D$  w.r.t the  $\alpha_i$ , subject to:

$$\sum_i^N \alpha_i y_i = 0, \quad 0 \leq \alpha_i \leq \gamma$$

- ▶ The solution must satisfy the conditions:

$$\mathbf{w} = \sum_i^N \alpha_i y_i \mathbf{x}_i \tag{6}$$

$$\sum_i^N \alpha_i y_i = 0 \tag{7}$$

$$\alpha_i = \gamma - \mu_i, \quad \forall i \tag{8}$$

$$\alpha_i [y_i (\mathbf{w}^\top \mathbf{x}_i + b) - (1 - \xi_i)] = 0, \quad \forall i \tag{9}$$

$$\mu_i \xi_i = 0 \quad \forall i \tag{10}$$

$$y_i (\mathbf{w}^\top \mathbf{x}_i + b) - (1 - \xi_i) \geq 0 \quad \forall i \tag{11}$$

# Lagrange Multipliers Solution

- ▶ From Eq.6, the solution is  $\mathbf{w} = \sum_i^N \alpha_i y_i \mathbf{x}_i$ . From Eq. 9,  $\alpha_i > 0$  when constraint 11 are exactly met, i.e.,  
 $y_i(\mathbf{w}^\top \mathbf{x}_i + b) - (1 - \xi_i) = 0$
- ▶ The points  $(\mathbf{x}_i)$  with  $\alpha_i > 0$  are the SUPPORT VECTORS.
- ▶ We have two kinds of support vectors:
  - ▶ Those for which  $\xi_i = 0$ : they lie on the edge of the margin.  
From Eqs. 10 and 8:  $0 < \alpha_i < \gamma$
  - ▶ Those for which  $\xi_i > 0$ : they have  $\alpha_i = \gamma$  and they lie on the wrong side of their margin.
- ▶ To estimate  $b$ , we can use Eq. 9 with any of the support vectors with  $\xi_i = 0$ .
- ▶ Once we have  $\mathbf{w}$  and  $b$ , the decision function can be written:

$$\hat{y} = \text{sign}(f(\mathbf{x})) = \text{sign}(\mathbf{w}^\top \mathbf{x} + b)$$

- ▶ The tuning parameter of this procedure is  $\gamma$ . Its optimal value can be estimated via cross validation.

# Optimization

- ▶ A constrained optimization problem over  $\mathbf{w}$  and  $\xi$

$$\min_{\mathbf{w}, b} \frac{\|\mathbf{w}\|^2}{2} + \gamma \sum_{i=1}^N \xi_i \quad (12)$$

subject to:  $y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 - \xi_i \quad i = 1, \dots, N$

(13)

- ▶ The constraint can be written more concisely as:

$$y_i f(\mathbf{x}_i) \geq 1 - \xi_i \quad (14)$$

together with  $\xi_i > 0$  is equivalent to

$$\xi = \max(0, 1 - y_i f(\mathbf{x}_i)) \quad (15)$$

- ▶ Hence the learning problem is equivalent to the unconstrained optimization problem over  $\mathbf{w}$ :

$$\min_{\mathbf{w}, b} \frac{\|\mathbf{w}\|^2}{2} + \gamma \sum_{i=1}^N \max(0, 1 - y_i f(\mathbf{x}_i)) \quad (16)$$

# Loss function

$$\min_{\mathbf{w}, b} \frac{\|\mathbf{w}\|^2}{2} + \gamma \sum_{i=1}^N \underbrace{\max(0, 1 - y_i f(\mathbf{x}_i))}_{\text{Hinge loss}} \quad (17)$$

- ▶  $y_i f(\mathbf{x}_i) > 1$ : points outside margin. No contribution to loss
- ▶  $y_i f(\mathbf{x}_i) = 1$ : points on margin. No contribution to loss (hard margin case)
- ▶  $y_i f(\mathbf{x}_i) < 1$ : points violates margin constraints. Contribute to loss.

# Hinge Loss




Figure: Hinge loss vs 0-1 loss

An approximation to the 0-1 loss. Convex?

# Implementation

- ▶ Solving the Quadratic Programming Problems (slow)
- ▶ Use an interior point method that uses Newton-like iterations to find a solution of the Karush–Kuhn–Tucker conditions of the primal and dual problems
- ▶ Platt's sequential minimal optimization (SMO) algorithm
- ▶ Stochastic **sub**-gradient descent algorithms.

# Non-linear SVM

- ▶ How can the above methods be generalized to the case where the decision function is not a linear function of the data?
- ▶ It turns out that the generalization to a nonlinear boundary can be accomplished in a straightforward way using a simple mathematical trick!
- ▶ One major observation (look at the dual objective function obtained earlier):

$$L_D = \sum_i \alpha_i - \frac{1}{2} \sum_i \sum_j \alpha_i \alpha_j y_i y_j \mathbf{x}_i \mathbf{x}_j \quad (18)$$

$$= \sum_i \alpha_i - \frac{1}{2} \sum_i \sum_j \alpha_i \alpha_j y_i y_j \underbrace{\langle \mathbf{x}_i, \mathbf{x}_j \rangle}_{\text{dot product}} \quad (19)$$

- ▶ The only way in which the data appear in the training problem is in the form of dot products.

# Non-linear SVM

- ▶ How about the solution function?
- ▶ From  $\mathbf{w} = \sum_i^N \alpha_i y_i \mathbf{x}_i$ , the solution function can be written as:

$$\begin{aligned}f(\mathbf{x}) &= \mathbf{w}^\top \mathbf{x} + b \\&= \sum_i^{N_s} \alpha_i y_i \mathbf{x}_i^\top \mathbf{x} + b \\&= \sum_i^{N_s} \alpha_i y_i \langle \mathbf{x}_i, \mathbf{x} \rangle + b\end{aligned}$$

where  $N_s$  is the number of support vectors.

⇒ Also in the solution function, the data appear in the form of dot products where the  $(\mathbf{x}_i)$ s are the support vectors.

# Non-linear SVM

- ▶ Now, suppose we first map the data to some high dimension Euclidean space using a mapping  $\Phi$ :

$$\Phi : \mathbb{R}^d \rightarrow \mathbb{R}^h$$

usually  $h > d$

- ▶ The idea is to enlarge the input space to achieve better training class separation.
- ▶ In general, linear boundaries in the **enlarged space** translate to **nonlinear boundaries** in the original space (true for any nonlinear mapping  $\Phi$ )

# Mapping

- ▶ Then, we compute the largest margin hyperplane in the new space  $\mathbb{R}^h$ .
- ▶ Of course, the training algorithm would only depend on the data through dot products in  $\mathbb{R}^h$ , i.e.,  $\langle \Phi(\mathbf{x}_i), \Phi(\mathbf{x}_j) \rangle$  where  $\Phi(\mathbf{x}_i) \in \mathbb{R}^h$ .
- ▶ Suppose we have a function (called **kernel function**)  $K$  that computes such dot products in the transformed space:

$$K(\mathbf{x}_i, \mathbf{x}_j) = \langle \Phi(\mathbf{x}_i), \Phi(\mathbf{x}_j) \rangle$$

- ▶ Then: all we need in the training algorithm is  $K$ , and we would never need to explicitly even know what  $\Phi$  is.
- ▶ Resulting procedure: replace  $\langle \Phi(\mathbf{x}_i), \Phi(\mathbf{x}_j) \rangle$  with  $K(\mathbf{x}_i, \mathbf{x}_j)$  everywhere in the training algorithm.
- ▶ Example of  $K$ :  $K(\mathbf{x}_i, \mathbf{x}_j) = e^{\frac{-||\mathbf{x}_i - \mathbf{x}_j||^2}{2\sigma^2}}$  (Gaussian kernel)

# Mapping

- ▶ The algorithm constructs a **linear** support vector machine in  $\mathbb{R}^h$ .
- ▶ It achieves this objective in roughly the same amount of time it would take to train on the un-mapped data.
- ▶ How can we use such a machine? In test phase, given the test points  $\mathbf{x}$ :

$$\begin{aligned}f(\mathbf{x}) &= \sum_i^{N_s} \alpha_i y_i \langle \mathbf{x}_i, \mathbf{x} \rangle + b \\&= \sum_i^{N_s} \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}) + b\end{aligned}$$

where  $\mathbf{x}_i$  are the support vectors and  $N_s$  is the number of support vectors.

- ▶ The fact that, through the kernel function  $K$ , we can work with vectors in input space, without even knowing the mapping function  $\Phi$  is known as the “**kernel trick**”

## Example - kernel functions

Example: an allowed kernel for which we can construct the mapping  $\Phi$

- ▶ Training data are vectors in  $\mathbb{R}^2$ .
- ▶ Suppose we choose  $K(\mathbf{x}_i, \mathbf{x}_j) = (\langle \mathbf{x}_i, \mathbf{x}_j \rangle)^2$ .
- ▶ We can find a mapping

$$\Phi : \mathbb{R}^2 \rightarrow \mathbb{R}^h$$

such that  $(\langle \mathbf{x}_i, \mathbf{x}_j \rangle)^2 = \langle \Phi(\mathbf{x}_i), \Phi(\mathbf{x}_j) \rangle$

- ▶ One such mapping is:

$$\Phi : \mathbb{R}^2 \rightarrow \mathbb{R}^3$$

defined as

$$\Phi(\mathbf{x}) = \begin{pmatrix} x_1^2 \\ \sqrt{2}x_1x_2 \\ x_2^2 \end{pmatrix}$$

where  $\mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$

## Example - kernel functions

- We can verify that this is indeed the case:

$$\begin{aligned} K(\mathbf{x}, \mathbf{y}) &= (\mathbf{x}, \mathbf{y})^2 = (x_1 y_1 + x_2 y_2)^2 \\ &= x_1^2 y_1^2 + x_2^2 y_2^2 + 2x_1 y_1 x_2 y_2 \end{aligned}$$

$$\begin{aligned} <\Phi(\mathbf{x}), \Phi(\mathbf{y})> &= \Phi(\mathbf{x})^\top \Phi(\mathbf{y}) \\ &= (x_1^2, \sqrt{2}x_1 x_2, x_2^2) \begin{pmatrix} y_1^2 \\ \sqrt{2}y_1 y_2 \\ y_2^2 \end{pmatrix} \\ &= x_1^2 y_1^2 + x_2^2 y_2^2 + 2x_1 y_1 x_2 y_2 \end{aligned}$$

- Note: in general neither the mapping  $\Phi$  nor the space  $\mathbb{R}^h$  are unique for a given kernel.

## Exercise - kernel functions

- ▶ You can verify the following two mapping  $\Phi$  also satisfy  $K(\mathbf{x}, \mathbf{y}) = \langle \Phi(\mathbf{x}), \Phi(\mathbf{y}) \rangle$  for kernel given above.
- ▶ Example 1

$$\Phi : \mathbb{R}^2 \rightarrow \mathbb{R}^3$$

$$\Phi(\mathbf{x}) = \frac{1}{\sqrt{2}} \begin{pmatrix} (x_1^2 - x_2^2) \\ 2x_1x_2 \\ (x_1^2 + x_2^2) \end{pmatrix}$$

- ▶ Example 2

$$\Phi : \mathbb{R}^2 \rightarrow \mathbb{R}^4$$

$$\Phi(\mathbf{x}) = \begin{pmatrix} x_1^2 \\ x_1x_2 \\ x_1x_2 \\ x_2^2 \end{pmatrix}$$

## Which functions are allowable as kernels?

- ▶ The mathematical properties that a function  $K$  must have so that a mapping  $\Phi$  and an expansion  $K(\mathbf{x}, \mathbf{y}) = \langle \Phi(\mathbf{x}), \Phi(\mathbf{y}) \rangle$  exist, have been studied and are well known (Mercer theorem)
- ▶ Two popular choices for  $K$  are:
  - ▶  $d^{th}$  degree polynomial:  $K(\mathbf{x}, \mathbf{y}) = (1 + \langle \mathbf{x}, \mathbf{y} \rangle)^d$
  - ▶ Radial basis:  $K(\mathbf{x}, \mathbf{y}) = e^{\frac{-||\mathbf{x}_i - \mathbf{x}_j||^2}{2\sigma^2}}$

# Importance of Parameter $\gamma$ for Non-linear SVMs

- ▶ Perfect separation of training data can in general be achieved in the enlarged space  $\mathbb{R}^h$ .
- ▶ Such perfect separation leads to the danger of overfitting the data. As a consequence, the classifier will generalize poorly.
- ▶ A proper setting of  $\gamma$  allows to avoid overfitting.
- ▶ Lets look at the objective function minimized by an SVM:  
$$\frac{1}{2}\|\mathbf{w}\|^2 + \gamma \sum_i^N \xi_i$$
 where  $\gamma$  is the penalty factor for errors.
- ▶ Large  $\gamma \rightarrow$  discourage any positive  $\xi_i \rightarrow$  tendency to overfit the data  $\rightarrow$  highly complicated decision boundary in input space.
- ▶ Small  $\gamma \rightarrow$  encourage a small value of  $\|\mathbf{w}\| \rightarrow$  larger margin  $\rightarrow$  more data on the wrong side of their margin  $\rightarrow$  smoother decision boundary in input space.
- ▶ In practice: we need to tune  $\gamma$  so to achieve best test error performance.

# Limitations and Open Issues

- ▶ Choice of the kernel
  - ▶ performance may depend on the chosen kernel, and on the values of associated parameters;
  - ▶ the best choice of a kernel for a given problem is still a research issue; (e.g., Latent Semantic Kernel for document classification)
- ▶ Speed and size for training and testing
  - ▶ Training: the evaluation of the dual objective function requires the computation of all dot products  $\langle \mathbf{x}_i, \mathbf{x}_j \rangle \Rightarrow$  time complexity  $O(N^2d)$  where  $N$  is the number of training data and  $d$  is the dimension.
  - ▶ Testing: need to evaluate:  $f(\mathbf{x}) = \sum_i^{N_s} \alpha_i y_i \mathbf{x}_i^\top \mathbf{x} + b$ .
    - ▶  $\Rightarrow$  time complexity  $O(M \cdot N_s)$  where  $M$  is the number of operation required to evaluate the kernel.
    - ▶  $\Rightarrow$  time complexity  $O(d \cdot N_s)$  when  $K$  is a RBF kernel,  $M = O(d)$ .
- ▶ Extension to multiple classes

## Acknowledgements

Slides adapted from Dr. Carlotta Domeniconi's *Pattern Recognition* at George Mason University  
and Notes on SVM by Andrew Ng